

CollectNaija

AI-Powered Collections & Payment Management Platform

“Collect what you're owed. Stay in control.”

Product Requirements & Software Development Blueprint

Version 1.0 • September 2026 • Confidential • For Founders, Product, Engineering & Investors

CORE FLOW

Customer !' Invoice !' Debt !' Payment !'
Reminder !' Conversation !' Collection !'
Receipt !' Analytics

AI PRINCIPLE

Professional • Respectful • Helpful • AI-assisted, not AI-controlled

STACK

React + Django + DRF + PostgreSQL +
Redis + Celery • Provider-agnostic AI

DOCUMENT CONTROL

Classification: Confidential — Internal & Selected Investors
Prepared by: CollectNaija Product Architecture Team (PM, Design, Engineering, AI, Security, Compliance)
Review Status: Draft v1.0 — pending founder validation & pilot feedback
Distribution: Founders • CTO • Lead Engineers • Design • AI Safety Reviewer • Legal/Compliance

This blueprint distinguishes FACT / REQUIREMENT / ASSUMPTION / HYPOTHESIS / RECOMMENDATION / FUTURE IDEA throughout. Do not present assumptions as proven market facts.

SECTION 00

How to Use This Blueprint

Read order, conventions, and acceptance gate before build

Who this is for

- Founders & Product: scope, roadmap, pricing, validation — sections 1–9, 45–54, 63.
- Engineering (FE/BE/Mobile/DevOps/QA): architecture, DB, API, security, testing — sections 15–44, 55–58.
- AI Team: agent, voice, safety, channel engine, cost — sections 10–14, 22–23, 30, 34, 36.
- Compliance/Legal & GTM: privacy, risk, compliance, metrics — sections 31, 43–51.

Conventions

Tag	Meaning	Example
FACT	Verifiable or evidence-backed	Nigeria has >100M mobile subscribers (NCC data)
REQUIREMENT	Must-have for launch/quality gate	Backend is source of truth for balances
ASSUMPTION	Working belief to validate	SMBs will pay ~5k–15k/mo for automation
HYPOTHESIS	Testable, may be wrong	WhatsApp outperforms SMS for collections
RECOMMENDATION	Suggested approach, not mandatory	Use Paystack abstraction first
FUTURE IDEA	Post-stabilization	Cash-flow forecasting & churn signals

Every financially significant action requires deterministic backend validation. The AI never fabricates balances, deadlines, penalties, or payment confirmations.

Golden Rule — AI-assisted, not AI-controlled

The AI is a respectful assistant. The business owner controls policies. The backend is the source of truth. Humans handle complex, disputed, or sensitive cases. Every page of this blueprint enforces that separation.

Contents

65 sections • ~70 pages • Architecture diagrams, ERD, state machines, API examples, checklists

1.	Executive Summary	
2.	Product Vision	
3.	Problem Definition	
4.	Market Validation Strategy	
5.	Target Customers	
6.	Personas	
7.	User Journeys	
8.	Product Requirements	
9.	MVP Scope	
10.	AI Collections Agent	
11.	AI Conversation Architecture	
12.	AI Voice Architecture	
13.	Communication Architecture	
14.	Collection Automation	
15.	Payment Architecture	
16.	Database Architecture	
17.	Entity-Relationship Diagram (ERD)	
18.	Multi-Tenant Architecture	
19.	Backend Architecture	
20.	Frontend Architecture	
21.	Mobile Architecture	
22.	AI Architecture	
23.	AI Safety Architecture	
24.	Security Architecture	
25.	Privacy Architecture	
26.	API Design	
27.	Webhook Architecture	
28.	Notification System	
29.	Reminder Engine	
30.	Human Escalation System	
31.	UI/UX System	
32.	Accessibility	
33.	Internationalization	
	Performance	
34.	Offline / Poor-Network Strategy	
35.	Testing Strategy	
36.	AI Testing	
37.	DevOps	
38.	Docker	
39.	CI/CD	
40.	Monitoring	
41.	Logging	
42.	Backups	
43.	Disaster Recovery	
44.	Monetization	
45.	SaaS Pricing Strategy	
46.	Unit Economics	
47.	Customer Acquisition	
48.	Metrics	
49.	Risk Register	
50.	Compliance Considerations	
51.	Build-vs-Buy Analysis	
52.	Technology Decisions	
53.	12-Month Roadmap	
54.	Weekly Development Process	
55.	Definition of Done	
56.	Git-Hub Strategy	
57.	Production Deployment Checklist	
58.	Pilot Launch Strategy	
59.	Future AI Roadmap	
60.	Flutter Roadmap	
61.	Global Expansion Strategy	
62.	Master Development Checklist	
63.	Glossary	
64.	Final Product Principles	
65.		

Reading path for builders: skim sections 1–9, then deep-dive 14–23 (automation & AI), 15–19 (data & backend), 26–30 (integration), 24–25/38–44 (hardening), 54–59 (execution). Mark every HYPOTHESIS for pilot validation before coding it into policy defaults.

SECTION 01

Executive Summary

Why CollectNaija, what we build, and what commercial outcome we chase

CollectNaija is a Nigeria-first, globally scalable SaaS platform for tracking customers, invoices, debts, payments and collections, with an intelligent AI Collections Agent that communicates respectfully via WhatsApp, voice, SMS and email. The product promise is simple: know who owes you, know how much, know when they promised to pay — and follow up automatically without damaging customer relationships.

- Problem: SMBs track debts in notebooks/Excel/WhatsApp/memory !' forgotten debts, late payments, disputes, wasted staff time, no history or accountability.
- Solution: Centralized flow Customer !' Invoice !' Debt !' Payment !' Reminder !' Conversation !' Collection !' Receipt !' Analytics, with policy-driven automation and an AI assistant that behaves like a professional, calm, business-aware human — not an aggressive collector.
- Commercial lens: Nigeria's credit-informal economy + WhatsApp ubiquity + mobile-money/bank-transfer rails make respectful, automated collections a high-frequency, high-ROI workflow. Land with SMBs (schools, clinics, traders, agencies), expand to branches and enterprise, monetize automation & AI usage separately.
- Architecture: React frontend, Django+DRF backend, PostgreSQL, Redis/Celery, object storage, provider-agnostic AI abstraction, channel abstraction (WhatsApp/SMS/Email/Voice), payment-provider abstraction, Docker/CI/CD on cloud.
- Guardrails: Backend is source of truth for all financial state; AI is tool-constrained, policy-gated, tenant-isolated, logged, and escalates early. Payment verification is webhook + idempotency, never frontend confirmation. Reminders cancel deterministically after verified payment.

FACT

Nigeria is a mobile-first market with WhatsApp as the dominant SMB-to-customer channel. (NCC & GSMA mobile penetration reports; observable SMB behavior).

HYPOTHESIS

Businesses will pay a premium for AI that reduces manual follow-up while preserving customer goodwill. Must validate via 20+ interviews and a paid pilot before pricing locks.

Pillar

What success looks like at 12 months

Paid SMBs

200–500 paying orgs, <8% monthly churn, 40% from WhatsApp-first verticals

Collection lift

15–30% reduction in average days-to-payment vs baseline

AI adoption

60% of paid orgs enable at least one AI channel; escalation <18%

Reliability

99.5% API availability, zero cross-tenant leaks in audit

What makes this investable

Not 'a chatbot that asks for money' — a financial collections operating system where automation compounds. Network of payment history + promise-to-pay + channel preference becomes defensible data moat per org, and anonymized benchmarks per vertical.

SECTION 02

Product Vision

“Collect what you’re owed. Stay in control.” — respectful, reliable, global-ready

CollectNaija helps businesses get paid without destroying the relationship between business and customer. The vision is a collections operating system that is simple for owners, powerful for collection teams, respectful to customers, reliable for financial records, safe for AI communication, and scalable internationally.

- North star: Every business knows its outstanding in one screen, every customer has a clear, respectful payment experience, every interaction is recorded, every naira is reconciled.
- Experience principles: Simplicity over configurability at first; clarity over cleverness; transparency over manipulation; consent and escalation over persistence.
- Nigeria-first doesn't mean Nigeria-only: localize currency, payment rails, language, time zone, and compliance behind abstractions; default brand remains Nigerian-credible and globally legible.
- AI posture: The agent is a professional assistant — concise, empathetic, context-aware, never threatening, never fabricating, always disclosing automation where required.

Vision hierarchy:

Purpose: Help businesses stay in control of money owed.

Positioning: Collect what you're owed. Stay in control.

Personality: Professional + Respectful + Calm + Helpful + Human-like + Concise + Empathetic + Business-aware

REQUIREMENT The AI must never threaten, insult, harass, publicly embarrass, or invent penalties/consequences. All wording comes from business-approved templates or policy-gated LLM generation with safety review.

RECOMMENDATION Keep the marketing promise narrow for year one: debt tracking + respectful automation. Defer 'predictive finance' claims until data moat exists.

SECTION 03

Problem Definition

From notebooks and screenshots to missed follow-ups and poor cash-flow visibility

Target businesses manage debts informally: notebooks, Excel, WhatsApp messages, paper invoices, bank screenshots, mental math, and ad-hoc calls. This creates systemic failure modes across the Cash !' Conversation !' Collection chain.

Current behavior	Failure mode	Business impact
Notebook/Excel	Forgotten debts, duplicate records	Revenue leakage, disputes
WhatsApp threads	No central history, staff-dependent memory	No accountability, knowledge loss on turnover
Manual calls	Inconsistent tone, wasted hours	Staff burnout, unprofessional impression
Screenshots for proof	Difficult reconciliation	Payment disputes, trust erosion
Memory-based promises	Missed follow-ups	Cash-flow unpredictability

Root causes

- No single source of truth for Customer !' Invoice !' Payment.
- No policy engine for when/how to remind — every reminder is a judgment call.
- No structured promise-to-pay — commitments live in chat history.
- No closed loop between payment verification and reminder cancellation !' customers reminded after paying.
- No tenant-isolated audit trail — who changed what, when, why is untraceable.

Opportunity

Centralization + automation + respectful AI turns a daily chore into a system: balances are authoritative, reminders are policy-driven, conversations are recorded, receipts are automatic, and analytics answer 'who owes what, when, and what to do next' in one glance.

FACT

SMB cash-flow visibility is low when receivables are unmanaged; late payments are a leading cause of working-capital stress for SMEs (World Bank SME finance literature).

ASSUMPTION Businesses that self-report 'we use WhatsApp to chase payments' are reachable and willing to trial a dedicated collections workspace — to be validated in Phase 0.

SECTION 04

Market Validation Strategy

Interview 20+ businesses before locking price or policy defaults

Phase 0 is discovery, not coding. The goal is to recruit pilot customers, measure current state, and kill attractive fictions before they become hard-coded flows.

Validation plan (3–4 weeks)

- Recruit 20–30 businesses across 5 verticals: private schools, pharmacies/clinics, trade/retail, agencies/freelancers, cooperatives/lending-adjacent (non-regulated). Include Lagos + one secondary city (Ibadan/Abeokuta) to test network/market variance.
- Interview guide (45 min): Show current tracking (photo of notebook/sheet), walk through last 5 collections, count outstanding, days-to-pay, channel used, time spent, tools tried, willingness to pay (Van Westendorp).
- Artifact per interview: one-page canvas — org size, monthly receivables, % overdue >7/30 days, primary channel, biggest pain, price sensitivity, pilot interest (yes/maybe/no).
- Success gate: "e12 'pilot-ready' orgs (active receivables, WhatsApp-reachable customers, willing to import 20+ customers/invoices). Otherwise narrow ICP or reposition.

Signal	How measured	Kill/Proceed threshold
Problem intensity	Hours/week on follow-ups, % overdue	"e6 hrs/week or "e20% overdue
Channel fit	% customers reachable on WhatsApp/SMS	"e60% reachable
Willingness to pay	Van Westendorp + 'pay today' test	"e40% say "e 5k/mo
Data readiness	Can provide customer/invoice list	"e70% can provide within 48h

Pilot design

- 10 pilot orgs, 4 weeks, white-glove import, weekly check-in, Slack/WhatsApp group. Track: days-to-payment, reminder response rate, promise-to-pay fulfillment, staff time saved, NPS, 'would pay' question.
- Pricing hypothesis tested in pilot: Free!'Starter!'Business!'Professional ladder, with AI minutes/messages metered separately. Do not lock annual discount until churn data exists.

REQUIREMENT No AI voice launch before pilot proves WhatsApp+SMS+Email baseline lifts collections. Voice is Phase 4, not Phase 1.

Anti-pattern
Do not build for 'all businesses'. Anchor on one wedge vertical first — schools (recurring fees, parent contact, term cycles) are the sharpest pilot fit for Nigeria.

SECTION 05

Target Customers

Start narrow — schools, clinics, traders — then expand to branches and enterprise

Initial customer profile (ICP) — Nigeria SMB that invoices repeatedly, has 30–1,000 active customers, collects via bank transfer or Paystack/Flutterwave, and already uses WhatsApp for follow-ups. Secondary profile: multi-branch SMEs needing role separation and audit.

Segment	Example	Receivables pattern	Why CollectNaija fits
Schools	Private nursery/primary	Term fees, per-student arrears	Recurring, parent contact, promise-to-pay
Clinics/Pharmacies	Health SMEs	Post-service billing	High dispute sensitivity !' respectful AI
Trade/Retail	Distributors	Credit sales	Volume, WhatsApp-first, receipt need
Services	Agencies, tutors	Milestone invoices	Professional tone, payment links
Branches	Multi-site SMB	Consolidated + branch view	RBAC, branch isolation, analytics

Non-targets (for now)

- Regulated lending/loan-recovery needing licensed collection practices — out of scope and high risk.
- Enterprises needing on-prem or deep ERP integration before product maturity — serve via API in Phase 5, not Phase 1.
- Businesses with <10 invoices/month and no follow-up pain — acquisition cost exceeds value.

ASSUMPTION Education vertical will have the shortest sales cycle due to term-driven urgency. Validate against traders who may have higher volume but lower willingness to formalize.

FACT Nigeria's SME sector is large but fragmented; trust and referrals outperform cold outreach. Pilot referrals are the primary acquisition lever in month 1–6.

SECTION 06

Personas

Owner, Staff, Customer — plus AI as a fourth actor with constrained agency

Primary personas

Persona	Goals	Pains	Success metric
Mama Emeka — Shop Owner	Know who owes, get paid faster	Notebook chaos, forgets to call	Days-to-payment !", no duplicate reminders
Mr. Ahmed — School Admin	Term fees collected, parent trust kept	Manual calls, disputes after payment	Collection rate !", complaints !"
Tolu — Collections Staff	Clear list, script, handoff	No history, tone inconsistency	Touches per collection !", escalation clarity
Customer — Parent/Buyer	Clear invoice, easy to pay	Confusing reminders, repeat asks after paying	One reminder !' pay !' receipt, no spam

AI as actor

• AI Collections Agent: monitors due dates, drafts reminders from approved templates + context, handles simple objections, records promises, escalates uncertainty. Never decides balance, status, or confirmation — queries backend.

- Voice Agent: identifies business, discloses automation properly, verifies identity before discussing sensitive amounts, records outcome, respects quiet hours and opt-outs.
- Policy Owner: Business owner defines reminder policy, tone, channels, limits, escalation. AI obeys policy + customer preference — never bypasses either.

REQUIREMENT Map each persona to RBAC roles (Owner, Admin, Staff, Read-only) with tenant-isolated data. Customer is not a platform user in v1 — they interact via payment link + comms channel without platform login.

Empathy guardrail


The person receiving a reminder is a customer, not an enemy. Every prompt, template, and voice script is reviewed for respect and clarity at a Primary-3 reading level where possible.

SECTION 07

User Journeys

End-to-end: Create !' Remind !' Respond !' Pay !' Verify !' Receipt !' Insight

Journey A — Business creates invoice and gets paid



- Entry:
Owner imports customers (CSV) or adds manually !' creates invoice with items, discount, tax !' total/ balance calculated on backend.
- Autom
ation:
Scheduler enqueues reminders per policy (e.g., 7d/2d before, due date, 2/7/14d after).
Channel engine picks WhatsApp/ SMS/Email per preference + history.
- Payme
nt:
Customer pays via link/transfer !' provider webhook !' id empotency !' balance updated !' pending reminders cancelled !' receipt sent !' AI notified.

- Edge: If customer responded 'already paid' before webhook, conversation is marked 'awaiting verification' and auto-resolves on webhook match or escalates after grace window.

Journey B — Customer responds to AI reminder

- AI sends polite, context-rich reminder (name, amount, due date, link, how to reach business). Customer replies in natural language.
- NLU classifies: Promise-to-Pay ('next week'), Already-Paid, Can't-Pay-Full, Stop-Contact, Confused, Wants-Human.
- AI handles with tools: records promise with date, requests reference if 'already paid', explains invoice from verified data, offers approved payment plan only if policy allows, pauses collection on opt-out, escalates on dispute/fraud/distress.
- Every turn is logged as AIMessage + CommunicationEvent; promise creates PromiseToPay; escalation creates SupportTicket + notifies staff.

Journey C — Staff handoff

- Trigger: dispute, repeated non-verifiable 'already paid', harassment/fraud flag, low confidence, or business escalation rule.
- System: creates ticket with full transcript, invoice snapshot, payment history, promised date, and suggested next step. Staff claims ticket, resolves, and policy engine pauses/resumes collection accordingly.
- Outcome: Customer experiences one continuous thread — human sees AI context, not a cold restart.

REQUIREMENT E2E journey must pass automated test weekly (section 39): business creates invoice !' reminder !' customer reply !' promise !' payment webhook !' idempotency !' cancellation !' receipt !' dashboard reflects truth.

SECTION 08

Product Requirements

Functional and non-functional requirements — the build contract

Functional — Core SaaS

- FR-01 Auth & org: signup, login, JWT, email verification, password reset, org creation, org switch, profile, last-login display, session management.
- FR-02 Customers: CRUD, search/filter/sort, outstanding/overdue aggregates, communication preferences, opt-out, CSV import, duplicate detection.
- FR-03 Invoices: create with items (qty, price), discount/tax, backend-computed totals, status (draft/sent/partial/paid/overdue/cancelled), due dates, customer scoping, audit log.
- FR-04 Payments: record (manual + provider), status (pending/successful/failed/refunded), idempotency key, receipt generation, balance propagation.
- FR-05 Reminders: template variables {{customer_name}}, {{amount}}, {{due_date}}, {{payment_link}}, {{business_name}}, {{invoice_number}}, channel/content preview, scheduled/sent/failed/cancelled states.
- FR-06 Collection Policy & Campaigns: per-org policy (frequency, channels, quiet hours, max attempts, grace, escalation, payment-plan rules) + per-cohort campaigns (e.g., School Fees).
- FR-07 AI Conversations: sessions, messages, actions, promise-to-pay, escalation, channel selection, safety layer, tenant-isolated tool calls.
- FR-08 Voice Calls: call attempts, duration, outcome, transcription, recording policy (metadata vs content separation), consent handling, retry with backoff.
- FR-09 Dashboards & Analytics: KPIs (outstanding, due today, overdue, collected this month), Needs Attention, cash-flow chart, AI activity timeline, insights (data-driven, never fabricated).
- FR-10 Admin & Compliance: RBAC, audit log, webhook logs, data retention controls, deletion/export (privacy).

Non-functional

NFR	Target	Rationale
Availability	99.5% monthly (Phase 1)	SMB collections cannot afford downtime at term end
Latency p95	API <400ms, AI turn <2.5s	Respectful turn-taking on WhatsApp/voice
Data isolation	Zero cross-tenant read	Security acceptance gate
Mobile perf	LCP <2.5s on 3G, bundle <200kb gz	Low-end Android, 3G reality
Cost guard	AI cost per org tracked & quota'd	Protect SaaS margins
Observability	AI/tool/comm/payment tracing	Debug revenue-critical flow

REQUIREMENT Every financial write (invoice total, payment status, balance) is backend-computed using Decimal (not float) with explicit rounding rules. Frontend math is preview-only.

SECTION 09
MVP Scope

What ships on Day 1 vs what waits — ruthless sequencing

MVP is ‘trustworthy debt tracking + respectful automation’ without voice. AI text assistant ships as controlled state machine with human escalation, but voice waits until text baseline is validated.

In	Out (Phase 2/3/4)
Auth, orgs, RBAC, customers, invoices, payments, receipts	Voice AI calls (Phase 4)
Reminders: WhatsApp + SMS + Email (provider abstraction)	Advanced call analytics & STT/TTS at scale
Policy engine: quiet hours, max attempts, channel prefs	Full campaign optimizer & prediction
AI text agent: reminder + simple conversation + promise + escalation	Multi-language NLU (Pidgin/Hausa/Yoruba/Igbo) — EN first
Payment verification: webhook + idempotency + auto-cancel	Complex payment plans (installments with interest)
Dashboard, reports, audit log, basic insights	Flutter mobile (Phase 6)

MVP acceptance checklist (ship blocker)

- Org can onboard in <5 min, import 20 customers, create 5 invoices, collect one payment end-to-end (manual + provider).
- Reminder respects quiet hours and max attempts; paid invoice cancels pending reminders within seconds of verified webhook.
- AI retrieves balance from backend, never invents amount, records promise-to-pay, escalates ‘already paid’ without proof, respects ‘stop contacting me’.
- Tenant isolation test suite green; audit log captures who changed what; receipts downloadable.
- P95 API <400ms in staging with seed of 10k invoices; offline banner and retry for poor network.

RECOMMENDATIONS

Cut pixel-perfect marketing site before you cut the webhook idempotency test. Revenue correctness beats polish in v1.

SECTION 10

AI Collections Agent

Professional, respectful, business-aware — the agent’s charter and toolbelt

The AI Collections Agent is a policy-gated, tool-constrained service that turns outstanding invoices into respectful conversations. It is not a free chatbot — it is a state machine that queries the backend for truth and acts only through audited tools.

Capabilities (and explicit non-capabilities)

- Can: monitor due dates, schedule reminders, send templated messages, answer invoice/balance questions from verified data, record promises, provide payment instructions/links, handle simple objections, escalate.
- Cannot: change invoice amounts, confirm payment without verification, waive fees without authorization, invent penalties or deadlines, reveal another tenant’s data, contact beyond policy limits, pretend to be human where disclosure is required.

Personality contract

You are the payment assistant for {business_name}. Tone: Professional, Calm, Respectful, Concise, Empathetic.
Always: identify the business, state the invoice/amount/due date from verified data, offer a payment link and a human path.
Never: threaten, insult, harass, invent consequences, or argue with 'I already paid.' Acknowledge, request reference, mark for verification, pause further if you do not know, say so and escalate. If distress/fraud/legal cue, escalate immediately.

Toolbelt (all audited, rate-limited, tenant-scoped)

Tool	Purpose	Side effect
get_customer / get_invoice / get_outstanding_balance	Verified read	None — read-only
get_payment_history / get_payment_link	Explain & pay	None
get_collection_policy / get_customer_preferences	Policy gate	None
record_customer_response / record_promise_to_pay	Capture intent	Creates PromiseToPay + follow-up
schedule_follow_up / send_payment_link	Next step	Enqueues CommunicationEvent
create_support_ticket / escalate_to_human	Handoff	Creates ticket, pauses AI turn
pause_collection / resume_collection	Consent	Updates policy state


REQUIREMENT Every tool call is authenticated (org+user/agent), authorized (RBAC + tenant check), validated (input/output schema), logged (AuditLog + AIAction), and rate-limited. High-impact writes (promise, pause, escalation) require backend deterministic validation.

SECTION 11

AI Conversation Architecture

State machine, memory, NLU, and generation with grounding

Conversation is a controlled state machine, not an open chat. States are persisted, transitions are validated, and context is bounded.



USTO
MER_R
EQUES
TED_S
TOP
cannot
go to R
EMIND
ER_SE
NT
without
explicit
resume
_collect
ion
after co
nsent.

**Memo
ry —
boun
ded
and g
overn
ed**

- R
etained
: custo
mer ide
ntity,
invoice
snapsh
ot, outs
tanding
, due
date, p
ayment
records
, last
90
days of
collecti
on inter
actions,
prefere
nces, pr
omised
date, a
pprove
d arran
gement
s,
ticket

refs.

- Not retained: full chat beyond retention window, biometric voice prints, unauthenticated personal data, other tenants' data (never in context).

- Retention: conversation content 12 months (configurable per org, lower where regulation demands) then soft-delete + hard-delete per schedule; metadata (who/when/channel/outcome) retained longer for

audit. Deletion is org-triggered and honoured within 30 days.

- Access: AI sees only the current org/customer/invoice slice via tools; humans see via RBAC + audit; no cross-tenant retrieval.

NLU Policy Generation pipeline

```
Inbound message ! normalize ! classify intent (Promise/AlreadyPaid/CannotPayFull/Stop/Confused/WantsHuman/Other) !
fetch verified context via tools ! policy gate (quiet hours, max attempts, tone, channel) !
generate grounding-blocked response (must cite tool outputs, no invention) ! safety layer (threat/harassment/finance-hallucination scan) !
act via tool(s) or send message ! log AIMessage + AIAction ! schedule next check
```

Grounding rule

The model may not output any amount, date, or status that did not come from a tool. A post-generation validator regex-matches currency/dates and cross-checks against the last tool payload; mismatch triggers a safe fallback ('Let me confirm that with the business') and an incident log.


FACT

State machines reduce hallucination and cost versus open-ended chat by constraining turns and tool order — essential for finance-adjacent AI.

SECTION 12

AI Voice Architecture

Approved calls, verification, STT !' LLM !' TTS, with strict retention separation



is a good time, listen (STT), handle intents as in text (promise/ already-paid/ escalate).

- 4) Record promise date, provide payment instructions, offer human transfer on request, end politely with summary SMS/ WhatsApp.
- 5) If not affirmed, wrong person, or distress/ fraud cue: end without sensitive disclosure, log INVALID _CONTACT or escalation.

Data handling

Data class	Content	Retention / access
Call metadata	status, duration, outcome, retry, consent flag	Retained 24 months, RBAC-visible
Recording/transcription	audio + text	Separate store, encrypted at rest, opt-in per org, 90-day default then delete, downloadable only by Owner/Admin with audit
Analytics	promise captured, escalation, call success	Aggregated, no PII sharing across tenants

REQUIREMENT Voice must not discuss balances with unverified persons. Verification is a gate, not a suggestion. Every dial, attempt, and outcome is a CommunicationEvent + VoiceCall + CallAttempt row for audit.


Nigeria reality

Voice is powerful for low-literacy or inbox-fatigued customers, but also sensitive. Default to WhatsApp/SMS first in v1; enable voice per customer only after text response pattern or explicit consent, and always respect time-of-day norms (e.g., 9am–7pm WAT configurable).

SECTION 13

Communication Architecture

Gateway + channel abstraction + delivery lifecycle



All outbound comm goes through a Communication Gateway that enforces policy, tenant isolation, templating, and lifecycle tracking. Channels are providers behind a common interface so WhatsApp (Meta Cloud API), SMS (Termii/ Africa's Talking), Email (SES/ SendGrid), and Voice (Twilio) can be swapped per org or per region without rewriting orchestration.

Message lifecycle

REQUESTED ! POLICY_GATED ! QUEUED ! SENDING ! SENT / FAILED ! DELIVERED? ! RESPONDED? ! BOUNCED? ! RETRY? ! CANCELLED (if paid/opt-out)
Every transition is a CommunicationEvent with provider message_id, timestamps, error codes, and cost.
Idempotency key = org_id + invoice_id + campaign_step + channel + scheduled_for (prevents duplicates on retries).

Template system

- Variables: {{customer_name}}, {{business_name}}, {{invoice_number}}, {{amount}}, {{balance}}, {{due_date}}, {{payment_link}}, {{support_phone}}. All rendered server-side with escaping; amounts formatted with Intl.NumberFormat and currency config.
- Tone presets (configurable): Formal, Friendly, Concise — mapped to template variants and LLM style hints, but policy still caps contacts.
- Language: template per locale; LLM generation is locale-aware but must still pass safety and grounding checks.

Delivery intelligence

- Gateway tracks per-channel delivery/read/reply rates per org/customer and suggests (but does not enforce) the best channel — final choice is always policy + customer preference.
- Fallback: if WhatsApp not delivered in N minutes, optionally retry via SMS if policy allows — never both simultaneously without explicit campaign step.

REQUIREMENT No direct provider call bypasses the gateway. All sends are queued via Celery with retry/backoff and respect the policy engine pre-send check. Failed sends surface to the owner with actionable reason, not raw provider codes.

SECTION 14

Collection Automation

Policy-driven scheduler that makes follow-up systematic

Automation is a deterministic scheduler (Celery Beat) + policy engine + campaign definitions. The AI generates wording, but the scheduler decides when and whether a message is allowed to go.

CollectionPolicy (per-organization) fields

Field	Type	Default (example)
reminder_intervals	list of offsets	-7d, -2d, 0, +2d, +7d, +14d
channels_per_step	map step!'channels	WhatsApp primary, SMS fallback
quiet_hours	time range + timezone	20:00–08:00 WAT, plus holidays
max_contacts_per_day/week	int	1/day, 3/week
min_interval_between_contacts	duration	24h
max_whatsapp/sms/calls	int	5 / 3 / 2 per campaign
grace_period	duration	48h after promise before next nudge
escalation_rules	predicates	dispute, already-paid-no-proof, distress
payment_plan_rules	bool + terms	allow_partial = false by default
tone / language	enum	Friendly, en-NG

CollectionCampaign (per cohort)

Example — School Fees Campaign: target = invoices where customer_type=student and due_date in term window. Steps: -7d WhatsApp, -2d WhatsApp, 0d WhatsApp+SMS, +3d WhatsApp, +7d AI call, +14d human review. Each org can clone and customize; steps are provider-agnostic and time-zone aware.

Scheduler loop (simplified)


```
@shared_task
def enqueue_due_reminders():
    for invoice in invoices_due_for_step(now()):
        policy = invoice.org.policy; prefs = invoice.customer.prefs
        if not policy_engine.allowed(invoice, policy, prefs, now()): continue
        with transaction.atomic():
            event = CommunicationEvent.create(..., idempotency_key=..., status='QUEUED')
            send_via_gateway.delay(event.id)
```

REQUIREMENT Policy engine executes before every outbound communication, including AI turns and voice dials. No path bypasses it. Policy changes apply to future enqueues only — already-QUEUED events keep their original policy snapshot for audit.

SECTION 15

Payment Architecture

Provider abstraction, webhook verification, idempotency, closed-loop cancellation



Principles

- Never trust frontend confirmation. Provider webhook +

verification is the only source of truth.

- Idempotency key = provider + provider_reference + amount + currency. Duplicate webhooks (retries) create exactly one Payment row and one balance mutation.

- Balance math uses Decimal with ROUND_HALF_UP per currency.

ncy
config
(NGN
has 0–2
decimals
; store
minor
units to
avoid
float).
Invoice
totals, ba
lances,
and
receipts
are
backend-
compute
d.

• Par
tial paym
ents
allowed;
overpay
ments
create
credit or
are
flagged
per org
setting;
refunds
are
explicit
and
reverse
balances
with
audit.

**Abstra
ction
layer**

Interface	Paystack example	Future provider
create_payment_link(amount, currency, reference)	initialize transaction	Flutterwave / Stripe adapter
verify_webhook(signature, payload)	HMAC check + retrieve transaction	Same
map_status(provider_status)	success/failed/pending	Same

refund(payment, reason)

refund API

Same

Webhook handler sketch

POST /api/v1/webhooks/payments/{provider} (no auth, HMAC verified, rate-limited, idempotent)
1) verify signature 2) parse event 3) fetch raw event -> WebhookEvent row (store before processing)
4) SELECT FOR UPDATE on idempotency_key 5) if exists, return 200 (ack) 6) else create Payment
7) update Invoice.balance inside transaction 8) emit: cancel pending reminders, send receipt, notify AI, update analytics
Always return 200 after durable store — never 500-retry a successfully stored event.

REQUIREMENT Payment confirmation via webhook must cancel pending reminders/calls within seconds and reflect in dashboard without manual refresh (poll or realtime). Customer must never receive a debt reminder after verified payment.

Reconciliation edge

Bank-transfer payments without provider link require manual 'Mark as received' with reference + proof upload, then a second-person approval if amount > threshold (configurable). Avoids social-engineering via fake screenshots.

SECTION 16

Database Architecture

Django, PostgreSQL, Redis — normalized core, auditable, soft-deletable, tenant-aware

Core tables are normalized; financial tables are append-friendly and audited. Soft deletion keeps history; hard deletion is a scheduled job after retention. Every tenant-scoped table has org_id and is filtered via Row-Level Security or manager scoping — never by frontend filter alone.

Group	Tables	Key constraints
Tenancy & Identity	Organization, User, Membership, Role	org slug unique, email unique, membership (org,user) unique
Business core	Customer, Invoice, InvoiceItem, Payment, PaymentMethod	invoice_number unique per org, Decimal amounts, FKs with RESTRICT
Collections	Reminder, CollectionCampaign, CollectionPolicy, CommunicationPreference, CommunicationEvent, Notification	policy one-to-one org, event idempotency unique
AI	AIConversation, AIMessage, AIAgentSession, AIAgentAction, PromiseToPay, PaymentPlan	conversation per customer/invoice, action audited
Voice	VoiceCall, CallAttempt	attempt per call, outcome enum, consent flag
Ops	SupportTicket, Escalation, AuditLog, WebhookEvent, Subscription	immutable audit, webhook payload retained

Cross-cutting concerns

- Soft delete: deleted_at timestamp + manager that excludes by default; unique indexes are partial (WHERE deleted_at IS NULL) to allow natural keys to be reused after deletion.
- AuditLog: actor_id, action, target_type/id, before/after JSON, ip, user_agent, created_at — never updated, only inserted.
- Indexes: (org_id, status, due_date) on Invoice; (org_id, customer_id, created_at) on Payment; (org_id, status, scheduled_for) on CommunicationEvent; GIN on

Customer search (name, phone, email).

- Data retention: messages/recordings/transcripts have org-configurable TTL; metadata retained longer; deletion jobs logged.

- Migrations: Django migrations with explicit SQL for partial indexes and check constraints (amount ≥ 0 , due_date not null when sent).

FACT

PostgreSQL + Django ORM is proven for multi-tenant SaaS with strong consistency needs. Soft-delete + partial unique indexes + RLS/manager scoping is a standard isolation pattern.

SECTION 17

Entity-Relationship Diagram (ERD)

Complete map — relationships, keys, and tenant spine

Tenant spine: Organization is the root for all business data. Every Customer, Invoice, Payment, Reminder, Policy, Conversation, Call, Ticket, and WebhookEvent belongs to one org. Users belong to orgs via Membership with a Role. This spine is enforced at DB and service layers.

```

Organization 1% % * User (via Membership)   Organization 1% % * Customer
Organization 1% % * Invoice                  Customer 1% % * Invoice
Invoice 1% % * InvoiceItem                   Invoice 1% % * Payment (via invoice_id) Payment *% % 1 Invoice
Invoice 1% % * Reminder                     Invoice 1% % * CommunicationEvent
Organization 1% % 1 CollectionPolicy         Organization 1% % * CollectionCampaign
Customer 1% % 1 CommunicationPreference     Customer 1% % * AIConversation
AIConversation 1% % * AIMessage              AIConversation 1% % * AIAction
AIConversation 1% % * PromiseToPay (0..1 active) AIConversation 1% % * PaymentPlan (optional)
AIConversation 1% % * VoiceCall              VoiceCall 1% % * CallAttempt
Organization 1% % * SupportTicket / Escalation Everything *% % * AuditLog (polymorphic)
WebhookEvent *% % 1 Organization (nullable until verified) Subscription 1% % 1 Organization
  
```

Representative DDL excerpts (illustrative)

```

organizations (id PK, slug UNIQUE, name, country, currency, timezone, created_at, deleted_at)
users (id PK, email UNIQUE, name, password_hash, created_at)
memberships (id PK, org_id FK!organizations RESTRICT, user_id FK!users CASCADE, role ENUM, UNIQUE(org_id,user_id))
customers (id PK, org_id FK, customer_code UNIQUE PER ORG WHERE deleted_at IS NULL, name, phone, email, outstanding DECIMAL, overdue DECIMAL, price_is_zero BOOLEAN, deleted_at)
invoices (id PK, org_id FK, customer_id FK, invoice_number UNIQUE PER ORG WHERE deleted_at IS NULL, status ENUM, due_date, subtotal, discounts DECIMAL, tax DECIMAL, balance DECIMAL, provider_ref VARCHAR, provider_ref UNIQUE, idempotency_key VARCHAR, created_at)
communication_events (id PK, org_id FK, invoice_id FK, customer_id FK, channel ENUM, template_id, status ENUM, provider_msg_id, idempotency_key VARCHAR, created_at)
ai_conversations (id PK, org_id FK, customer_id FK, invoice_id FK, state ENUM, channel, started_at, ended_at, escalated_at)
promise_to_pay (id PK, org_id FK, conversation_id FK, amount DECIMAL, promised_date DATE, status ENUM, fulfilled_at)
  
```

ERD visual (schematic)

Below is a simplified block view. In production, generate a full ERD from pgAdmin / Django schema with `django-extensions graph\_models` and keep it in /docs/erd.pdf versioned with migrations.



- All FKs are DEFERRABLE with RESTRICT on tenant root to prevent accidental cascade deletes; soft delete preserves history.
- Check constraints: amount >= 0, balance >= 0, total = subtotal - discount + tax (validated in model clean + DB check).
- GIN trigram indexes on customer name/phone for fast WhatsApp-era search; BRIN on created_at for time-series dashboards.

SECTION 18

Multi-Tenant Architecture

Isolation at DB, service, and API — never by frontend

Every request is tenant-scoped before any business logic. The org is resolved from JWT (user's active org) or API key, then enforced via manager scoping and, where needed, PostgreSQL Row-Level Security (RLS). Cross-tenant reads fail closed with 404 (not 403) to avoid enumeration.

Layers

Layer	Mechanism	Guarantee
Auth	JWT contains org_id + role, short-lived access + refresh	No org switch without re-auth
Middleware	CurrentOrg middleware sets request.org, thread-local for Celery	All downstream code reads request.org
ORM	TenantManager: .for_org(org) != filter(org_id=org.id)	Querysets without org raise
DB	RLS policies: USING (org_id = current_setting('app.org_id'))	Even raw SQL respects tenant
Storage	Object keys prefixed by org_id/	No object cross-read
Cache	Keys namespaced org:{id}:...	No cross-cache leak
Logs	All log lines include org_id, user_id	Forensic trace

Testing isolation (mandatory)

```
def test_cross_tenant_invoice_access_denied(client):
    orgA, orgB = create_orgs(2)
    inv = create_invoice(orgA)
    tokenB = login(orgB.owner)
    assert client.get(f'/api/v1/invoices/{inv.id}', headers=tokenB).status_code == 404
```

Branch model


- Branches are not tenants — they are sub-scopes within an org (branch_id nullable on customer/invoice). RLS remains org-level; branch filtering is service-level. Enterprise orgs can enforce branch RBAC (staff sees only branch invoices).
- Demo workspace is a real org flagged is_demo=true with isolated seed data and a visible badge — never mixed with production analytics.

REQUIREMENT Every query on a tenant table must include `org_id`. Code review and static analysis (custom lint) flag `raw .objects.filter` without `.for_org`. Tenant-isolation tests run in CI with seeded dual-org fixtures.

SECTION 19

Backend Architecture

Django + DRF, 12-factor, API-first, job-queue aware



Django app map

apps/tenancy, accounts, customers, invoices, payments, comms, collections, ai, voice, webhooks, audit, subscriptions
Each app: models.py, managers.py (tenant-scoped), serializers.py, views.py, urls.py, tasks.py, policies.py, tests/
Cross-app imports only via service layer (services.py) to avoid circular deps — e.g., collections service calls payments service for balance.

Request lifecycle

- NGINX !' Gunicorn !' Django middleware chain: RequestID !' SecurityHeaders !' CurrentOrg !' Auth !' RateLimit !' PolicyGate !' View !' AuditLog !' Response.
- All views are DRF ViewSets with explicit permission_classes (IsAuthenticated + HasOrgRole + TenantScoped). Serializers validate input; services enforce business rules; models enforce DB constraints.
- Heavy work (provider calls, PDF receipts, AI turns, comm sends) is enqueued — HTTP returns 202 with job id where appropriate; idempotency keys prevent double-enqueue.

API versioning & error contract

```
Base: /api/v1 (version in URL, never header-only)
Success: { success:true, data:{...}, meta:{page, per_page, total} }
Error: { success:false, message:'...', errors:{field:['...']}, code:'VALIDATION_ERROR' }
Codes: VALIDATION_ERROR, AUTH_REQUIRED, FORBIDDEN, NOT_FOUND, RATE_LIMITED, CONFLICT, PROVIDER_ERROR,
IDEMPOTENCY_REPLAY
```

REQUIREMENT All money math uses Decimal + currency config; FloatField is forbidden for amounts (lint rule). Every financial state change writes an AuditLog row in the same transaction.

SECTION 20

Frontend Architecture

React + TypeScript + Tailwind, mobile-first, offline-aware, generated-API ready

Frontend is a Vite + React 18 + TypeScript SPA (Next.js upgrade path reserved) with Tailwind, React Router, Recharts, lucide-react, and daysjs. It is already implemented as a production-grade shell: AppShell with desktop sidebar + bottom nav, protected routes, JWT handling, toast, and API-ready error contract.

Structure (as shipped)

```
src/components/ui Button, Input, Select, Badge, Card, Modal, Skeleton, EmptyState, Toast
src/layouts      AppShell (44px targets, sticky header, offline banner)
src/pages        Landing, Login/Signup, Onboarding(8-step), Dashboard, Customers(+Detail), Invoices(+Detail), Payments, Reminders, Reports, Settings
src/services     api.ts (apiFetch, JWT, pagination), mock.ts (labelled demo), store.ts (local demo state until API)
src/hooks/useAuth JWT, protected route gate      src/i18n keys (en-NG, NGN, Africa/Lagos)
```

Key UX decisions (spec-compliant)

- Mobile bottom nav: Home/Customers/Invoices/Payments/More — 44px min targets, no overflow, horizontal scroll only where needed.
- Cards on mobile !' tables on desktop; search/filter/sort with debounced query; modals preserve input, disable on submit, surface field errors from {errors} contract.
- Dashboard KPIs (Outstanding / Due Today / Overdue / Collected) are live from API (demo seed labelled); cash-flow AreaChart with 7/30/90/365 ranges; Needs Attention sorted most-overdue first.
- Invoices: items/discount/tax !' preview total on frontend, authoritative total from backend; Payments: method/status chips (pending/successful/failed/refunded), receipt drawer; Reminders: status/scheduled/sent/failed/cancelled with channel icons.
- Trust: Demo Workspace badge, last-login/sessions, audit log link, subscription transparency, analytics stubs without PII, offline banner via navigator.onLine.

Quality bars

- Performance: manualChunks (vendor/chart/ui), lazy routes, pagination, skeleton loaders, image optimization, gzip/br.
- A11y: semantic HTML, focus rings, keyboard nav, color+text+icon badges, reduced-motion support.
- i18n: Intl.NumberFormat for currency, date-fns/daysjs with timezone, locale keys — no hard-coded 'NGN' strings in components (use formatCurrency).
- Testing: Vitest + Playwright (critical E2E: create invoice !' reminder !' pay !' receipt !' dashboard). Contract tests against mock API.

RECOMMENDATIONS

Keep mock.ts but gate it behind VITE_USE MOCK and label every demo datum 'Demo — not real money' until API lands. Remove mock from production build.

SECTION 21

Mobile Architecture

Flutter — but only after SaaS validation (Phase 6)

Mobile is Flutter (see section 61 roadmap). In Phases 1–5 the web is responsive and PWA-capable; no native build until paid retention is proven. When built, Flutter shares the same Django API and RBAC; no separate business logic on device.

Concern	Web (now)	Flutter (later)
Auth	JWT in httpOnly cookie + memory	Secure storage, biometric opt-in
Offline	Banner + retry, queue payment intent	SQLite queue, background sync
Push	Web push (optional)	FCM/APNs for reminders/escalations
Biometrics	N/A	Face/Fingerprint for sensitive actions
Distribution	Vercel/host	Play Store + App Store (staged rollout)

REQUIREMENT Do not build Flutter before: "e100 paying orgs, <8% churn, and text AI proving value. Building mobile first is a known way to burn runway without retention signal.

PWA interim
Add PWA manifest + service worker for installability and offline queue of idempotent payment intents before Flutter ships — low cost, high trust on low-end Android.

SECTION 22

AI Architecture

Provider-agnostic LLM abstraction, cost-aware routing, deterministic grounding

Provider abstraction

```
class AIProvider(ABC):
    generate_response(prompt, tools, policy) -> LLMResponse
    classify_message(text) -> Intent
    extract_payment_intent(text) -> PromiseToPay | None
    summarize_conversation(messages) -> Summary
    determine_escalation(context) -> bool
    generate_voice_script(intent, context) -> Script
Providers: OpenAI, Anthropic, Google, local — swapped via config, no app rewrite.
```

Routing & cost control

Task	Model class	Why
Intent classify / promise extract	Small fast/cheap (e.g., nano)	High volume, low hallucination risk
Reminder drafting	Small–medium	Templated, policy-gated
Full conversation turn	Medium	Needs context + tool discipline
Escalation summarization	Medium–large	Needs nuance for human handoff

Context budget

- Per turn, context is at most: org policy summary + customer/invoice snapshot + last 8 turns + tool outputs. Older history is summarized, not replayed verbatim — caps tokens and leakage.
- System prompt is never included in customer-visible transcript; it is versioned and hashed for audit ('prompt version' logged per AIMessage).
- Token usage and cost per org are metered (AIAgentAction tokens + cost_minor) and surfaced in Settings ! Usage. Quotas enforce hard stops.

Grounding architecture

LLM output is not trusted for financial facts. A deterministic validator cross-checks every currency/date/status string against the last tool payload. Mismatch ! safe fallback + incident. Tool outputs are rendered into the prompt as a 'grounding block' that the model is instructed to cite.

REQUIREMENT AI provider switching must not change tenant isolation, safety rules, or grounding checks. Those live in the app layer, not the provider SDK.

SECTION 23


AI Safety Architecture

The Safety Layer that sits between the agent and the customer

Safety is not a prompt — it is a service that inspects inputs, outputs, tool calls, and contact policy before any customer-visible action. It is tested independently and versioned.

Threat	Defense	Verification
Threats/harassment language	Output classifier + blocklist + policy tone gate	Red-team set: 200 aggressive prompts !' 0 sends
Excessive contact	Policy engine pre-send + cooldown	Fuzz: 10k schedule simulations !' 0 over-limit
Reveal other tenant data	Tenant-scoped tools + output scan for org_id mismatch	Dual-org test !' 0 leaks
Prompt injection	Input sanitization + tool auth + instruction hierarchy	Injection corpus !' tools not called
Hallucinated balance	Grounding validator + tool-citation required	Balance fuzz !' 100% fallback on mismatch
Unauthorized payment confirm	No tool can set payment status to successful	Code search: 0 writes to Payment.status=success except webhook

Architecture



SMS threads are already addressed to the on-file contact, but messages avoid full account details unless the channel is verified and policy allows.

- The AI discloses automation where required by channel/region (configurable disclosure sentence).

- All blocks and fallbacks are logged as AI Agent Action with reason_code for review — silence is not an option; every blocked send has a human-readable alternative.

REQUIREMENT Safety tests are blocking in CI. Any change to prompts, tools, or policy defaults must re-run the AI safety suite (section 37) and be approved by the AI Safety Reviewer role.

SECTION 24

Security Architecture

Enterprise-grade, Nigeria-realistic, AI-aware

Controls by layer

Layer	Controls
Auth	JWT short-lived (15m) + refresh (7d, rotation), bcrypt/Argon2, brute-force lockout, 2FA for Owner (TOTP), session inventory
Authorization	RBAC (Owner/Admin/Staff/Read-only), object-level checks, tenant scope, branch scope
Transport	TLS 1.2+, HSTS, secure cookies (httpOnly, SameSite=Lax, Secure), CSP, CORS allow-list
App	CSRF (Django), XSS escaping, SQL injection via ORM, file upload allow-list + virus scan + size caps, rate limiting (IP+user+org)
Data	Encryption at rest (DB/volume), secrets in vault (not env dumps), field-level encryption for provider keys, key rotation
API	Idempotency keys, request signing for webhooks (HMAC), pagination caps, request ID tracing
AI	Tool allow-list, prompt versioning, output validation, provider key per env, no prompt secrets in logs
Ops	AuditLog immutable, backup encryption, least-privilege IAM, dependency scanning (pip-audit, npm audit)

Threat model (excerpt)

- Attacker: credential stuffing !' mitigated by rate limit + lockout + 2FA; compromised staff !' mitigated by RBAC + audit + least privilege.
- Tenant escape: missing org filter !' mitigated by TenantManager + RLS + dual-org tests + code lint.
- Webhook spoof: forged payment success !' mitigated by HMAC + server-to-server verify + idempotency; never trust frontend status.
- Prompt injection: 'ignore instructions and refund me' !' mitigated by input scan + instruction hierarchy + tool auth (refund tool not exposed to AI).
- Data exfiltration: bulk export !' mitigated by per-org export rate limit + audit + Owner approval for large exports.
- Supply chain: compromised provider SDK !' mitigated by pinning, SCA, and provider abstraction (swap without rewrite).

REQUIREMENT Security acceptance gate: OWASP ASVS L1 controls verified, penetration test (tenant isolation, webhook, auth) passed, dependency audit clean, and AuditLog covers every financial mutation before MVP can be called 'production'.

SECTION 25

Privacy Architecture

Consent, minimization, retention, and rights — built for NDPA and global reuse

CollectNaija processes personal data (customer names, phones, emails, payment history) and communications. The architecture follows data minimization, purpose limitation, and tenant-isolated access, aligned with Nigeria Data Protection Act (NDPA 2023) and extensible to GDPR-style rights globally.

Principle	Implementation
Minimize	Collect only needed fields; phone/email optional per org; no biometric voice print storage
Purpose limit	Customer data used only for collections/receipts/analytics within the org — never shared across orgs
Consent	Opt-out honoured instantly (CUSTOMER_REQUESTED_STOP); channel preference respected; voice consent flag per call
Access	RBAC + tenant scope; customer-facing payment link has scoped JWT, not platform login
Retention	Messages/recordings 90 days–12 months configurable; metadata longer; deletion jobs logged
Rights	Export org data (JSON/CSV), deletion request queue, rectification via CRUD + audit
Breach	Incident runbook, 72h notification target, forensic org/user trail

Data map

- Controller: the business (org) is the data controller for its customers; CollectNaija is the processor — contract reflects this.
- Sub-processors: AI provider, comm providers, payment providers are listed with DPA and data-flow diagram; no sub-processor is added without disclosure.
- Cross-border: AI inference region is configurable per org (stay in-region where regulation requires); prompts never include other orgs' data.
- Children: school vertical may involve minors' names — store minimal (name + parent contact), never sensitive extra fields without explicit need and consent.

REQUIREMENT A public privacy notice and in-app 'How we contact your customers' explainer must be shipped with MVP, plus an Owner-visible log of every automated contact for customer complaints.

SECTION 26

API Design

RESTful, versioned, tenant-scoped, paginated, idempotent

Resource map (v1)

```
POST /api/v1/auth/signup, /login, /refresh, /logout
GET/POST /api/v1/organizations, /organizations/{id}/switch
GET/POST /api/v1/customers, /customers/{id}, /customers/import
GET/POST /api/v1/invoices, /invoices/{id}, /invoices/{id}/items, /invoices/{id}/send
GET/POST /api/v1/payments, /payments/{id}/receipt, /payments/verify (internal)
GET/PUT /api/v1/policies (org policy), /campaigns, /campaigns/{id}/steps
GET/POST /api/v1/reminders, /communications, /communications/{id}/retry
GET/POST /api/v1/ai/conversations, /ai/conversations/{id}/messages, /ai/actions
GET/POST /api/v1/voice/calls, /voice/calls/{id}/attempts
GET/POST /api/v1/tickets, /tickets/{id}/messages, /webhooks/payments/{provider}
```

Example — create invoice

```
POST /api/v1/invoices Authorization: Bearer <jwt> Idempotency-Key: <uuid>
{
  "customer_id": "cus_...",
  "due_date": "2026-10-05",
  "items": [{"name": "Term fee", "qty": 1, "unit_price_minor": 8500000}],
  "discount_minor": 0, "tax_minor": 0,
  "currency": "NGN"
}
// 201 { success:true, data:{ id, number:'INV-2026-0042', total_minor, balance_minor, status:'sent' }}
// Duplicate Idempotency-Key ! 200 with original invoice (no double-create)
```

Conventions

- Pagination: ?page=1&per_page=20 (max 100), response meta {page, per_page, total, total_pages}.
- Filtering: ?status=overdue&due_before=2026-10-01&q=ahmed via ?q trigram search.
- Sorting: ?sort=due_date&order=asc; only allow-listed fields sortable.
- Idempotency: Idempotency-Key header on all POST/PUT that create or mutate money/comm state; server stores key+hash for 24h.
- Rate limiting: 60 req/min per user, 200/min per org, 10/min per IP for auth; headers X-RateLimit-Remaining.
- Errors use the standard contract with code and field errors; 404 hides existence across tenants.

REQUIREMENT Every list endpoint must be tenant-scoped and require org context; raw queryset access without .for_org is a CI failure. OpenAPI spec is generated from serializers and kept in /docs/openapi.yaml.

SECTION 27

Webhook Architecture

Durable, verified, retried, and replay-safe

Webhooks are the only way payments become 'successful' and the only way provider delivery receipts become 'delivered'. The handler is designed for at-least-once delivery from providers.

Storage first, processing second

```
1) NGINX !' Django: POST /api/v1/webhooks/payments/{provider}
2) Verify HMAC signature (constant-time), check timestamp skew (±5m)
3) Insert WebhookEvent { provider, event_id UNIQUE, payload JSONB, signature, received_at, processed:false } in transaction
4) Return 200 immediately (ack) — provider retries stop
5) Enqueue process_webhook_event.delay(event.id) — idempotent worker picks it up
6) Worker: SELECT FOR UPDATE on idempotency_key !' create Payment once !' update Invoice !' cancel reminders !' receipt !' AI notify
```

Reliability

- Provider retries (exponential) are expected; handler is idempotent — same event_id returns 200 without double-processing.
- DLQ: poison events (bad signature, malformed payload) go to WebhookEvent with error_code and are surfaced in Ops dashboard, not retried blindly.
- Replay: Owner/Admin can replay a stored event (creates a new processing attempt with same idempotency guard) for manual recovery.
- Outbound webhooks (future: org !' external ERP): HMAC-signed, retry with backoff, per-org endpoint config, delivery log with status/latency.

Security

- No auth cookie on webhook endpoints; only HMAC. IP allow-list optional but not relied upon (providers rotate IPs).
- Payload size cap (256KB), JSON schema validation, and rate limit per provider (100/min) to absorb floods.
- Secrets rotation: provider webhook secrets versioned; old+new accepted during rotation window, then old revoked.

FACT

At-least-once webhook delivery is the industry norm (Paystack, Stripe). Designing for exactly-once at the network layer is incorrect — idempotency at the application layer is the fix.

SECTION 28

Notification System

In-app, push-ready, preference-aware

Notifications tell the owner/staff what needs action without becoming another spam channel. They cover: new payments, failed comms, escalations, promises due, 'already paid' awaiting verification, quota warnings, and system incidents.

Notification	Trigger	Default recipient
Payment received	Webhook verified	Owner + Staff (invoice owner)
Payment failed	Provider status failed	Owner + Staff
Escalation created	AI or policy rule	Owner + assigned Staff
Promise due today	PromiseToPay promised_date = today	Owner + Staff
Already-paid unverified >48h	No matching Payment	Owner
Comm failed	Provider hard bounce	Owner
Quota 80%	AI messages/minutes threshold	Owner

Preferences & channels

- Per-user notification preference: in-app only vs in-app + email; per-org digest (daily 8am WAT) vs real-time.
- In-app inbox lives in AppShell (bell icon with dot), paginated, mark-read, deep-links to invoice/conversation/ticket.
- Future: mobile push via FCM; web push optional. No SMS to staff for routine events — reserve SMS for customer comms.

Implementation

Notification model: id, org_id, user_id, type, title, body, link, read_at, created_at
Fan-out via Celery task on domain events (payment_verified signal !' create notifications)
Realtime: polling (30s) in Phase 1, WebSocket (Django Channels) in Phase 3 if needed — not required for MVP

REQUIREMENT Every notification links to the authoritative object (invoice/payment/ticket) and shows the same balance/status as the backend — never a cached stale amount.

SECTION 29

Reminder Engine

Deterministic heartbeat that turns policy into scheduled sends

The reminder engine is a Celery Beat periodic task that scans for invoices crossing a policy step, creates `CommunicationEvents`, and enqueues gateway sends. It is the only component allowed to create reminders at scale — staff can create one-offs, but scale comes from the engine.

Tick (every 5 minutes)

```
for invoice in Invoice.objects.for_org(org).filter(status__in=['sent','partial','overdue'], due_date__is_not_null):
    step = policy.next_step_for(invoice, now()) # maps offset to step
    if step is None: continue
    if CommunicationEvent.exists(invoice, step, idempotency_key): continue
    if not policy_engine.allowed(invoice, policy, prefs, now()): continue # quiet hours, max attempts, cooldown
    CommunicationEvent.create(invoice, step, channel=engine.pick_channel(...), scheduled_for=now(), status='QUEUED')
    send_via_gateway.delay(event.id)
```

Cancellation

- On `payment_verified(invoice)`: UPDATE `communication_events` SET `status='CANCELLED'` WHERE `invoice_id=?` AND `status IN ('QUEUED','SCHEDULED')` AND `scheduled_for > now()` — single transaction with balance update.
- Cancellation is also triggered by opt-out (`CUSTOMER_REQUESTED_STOP`) and by human pause — same path, different `reason_code`.
- Cancelled events stay in the log with reason (`PAID`, `OPT_OUT`, `PAUSED`, `SUPERSEDED`) — the timeline shows 'Reminder cancelled — payment received' so the owner trusts the loop closed.

Failure handling

- SOFT fail (provider throttled): retry in 10m/30m/2h with backoff, preserve idempotency key, respect policy max on retry count.
- HARD fail (invalid phone): mark `FAILED`, surface 'Invalid contact' badge on customer, suggest correction, do not retry same channel blindly — try fallback only if policy allows.
- Monitoring: queue depth, scheduled vs sent lag, failure rate by channel, cancellation rate (high cancel rate after send = verification lag to fix).

REQUIREMENT Reminder scans are sharded by org and use `SELECT FOR UPDATE SKIP LOCKED` to allow multi-worker parallelism without double-enqueue. Every scan logs its window and count for audit.

SECTION 30

Human Escalation System

AI knows when to stop — staff get full context

Escalation is a first-class workflow, not an afterthought. The AI creates a SupportTicket with full transcript, invoice snapshot, payment history, promise history, and a recommended action. Staff claim, resolve, and the system pauses/resumes collection accordingly.

Triggers (any one != escalate)

- Customer disputes the debt / says invoice is wrong.
- Customer says 'I already paid' and no matching Payment found within grace window or reference invalid.
- Customer requests a human / says 'speak to someone'.
- Fraud/harassment/distress cue or legal language detected.
- AI confidence low, contradictory history, or business-defined threshold (e.g., balance > !500k and overdue >30d).
- Unusual arrangement request beyond approved payment-plan rules.
- Repeated fallback failures or invalid contact.

Ticket model

Field	Content
ticket_number	TKT-{org_slug}-{seq} (human-readable)
customer/invoice snapshot	JSON snapshot at escalation time (prevents drift)
transcript	Last N turns + tool outputs + safety flags
reason_code	DISPUTED ALREADY_PAID WANTS_HUMAN FRAUD DISTRESS LOW_CONFIDENCE POLICY_THRESHOLD
assignee/status	unassigned != claimed != in_progress != resolved / reopened
resolution	What staff did + whether collection resumed/paused

Handoff UX

- Staff sees ticket in Reminders/Reports queue with SLA (e.g., respond within 4h). One click opens CustomerDetail with conversation transcript pinned and 'Continue as human' composer that still logs via same channel.
- After resolution, staff chooses: Resume collection (with new policy), Pause until date, Mark disputed (pauses until admin override), or Write-off (with approval).
- Customer experiences continuity — human message appears in same WhatsApp thread; no 'please repeat your invoice number'.

REQUIREMENT Escalation is never silent: the customer receives a handoff message ('I'll connect you with the team'), the owner gets a notification, and the ticket SLA is tracked. AI does not attempt another turn after escalation until staff releases it.

SECTION 31

UI/UX System

Clean, trustworthy, mobile-first — Tailwind with semantic tokens

Design system is already in production in this repo: Inter font, brand #0F4C81, semantic tokens (emerald/amber/red), rounded-2xl cards, 44px targets, and a landing!onboarding!dashboard!detail hierarchy that an SMB owner can follow without training.

Token	Value	Usage
brand-600	#0F4C81	Primary actions, active nav
emerald	#059669	Success, paid, collected
amber	#D97706	Due soon, warning
red	#DC2626	Overdue, failed, danger
slate-50	#F8FAFC	Page background
radius	16–20px	Cards, buttons, inputs
type	Inter, 14px base	Headers 16–24px, body 14px, meta 12px

Information architecture

- Top nav: brand + search + notifications + user; sidebar: Home/ Customers/Invoices/Payments/Reminders/Reports/Settings; bottom nav mirrors on mobile.
- Dashboard: greeting + live KPIs + Needs Attention (most overdue first, View + Send Reminder per row) + cash-flow chart + recent invoices table.
- Detail pages: CustomerDetail shows outstanding/overdue + invoices + payments + reminder timeline; InvoiceDetail shows items + balance + receipt + comm history; Payments shows method/status/ receipt drawer.
- Empty/error/skeleton/zero states for every list — never a blank screen; every demo datum is labelled ‘Demo’ until replaced by API data.

Content & tone

- UI copy is simple English, no jargon: 'Add customer', 'Create invoice', 'Record payment', 'Send reminder'.
- Financial displays always pair color with text+icon (badge) for color-blind legibility; amounts use `formatCurrency` with currency code.
- Destructive actions (delete, write-off) require typed confirmation or second-person approval above threshold.

REQUIREMENT Every interactive element on mobile must be "e44px and no content may overflow viewport. Two consecutive pages with overflow in Playwright mobile viewport is a CI failure.

SECTION 32

Accessibility

WCAG 2.1 AA as baseline — not a backlog item

Accessibility is a product quality attribute, not a phase. The app targets WCAG 2.1 AA, with focus on keyboard, screen-reader, and color-contrast needs that are acute on low-end devices in bright sunlight.

- Semantic HTML: <nav>, <main>, <table> with <th scope>, <button> for actions, <a> for navigation — no div-buttons.
- Focus: visible ring (brand), logical tab order, skip-to-content link, modals trap focus and return it on close.
- Contrast: text "e4.5:1, large text "e3:1; badges use color+text+icon, charts use patterns/labels beyond color.
- Screen readers: aria-label on icon buttons, aria-live on toasts and KPI updates, table captions, form errors linked via aria-describedby.
- Motion: prefers-reduced-motion disables chart animation and modal transitions.
- Touch: 44px targets, 8px min spacing between tappable, no hover-only actions.

Testing

- Automated: axe-core in CI (0 critical violations), Lighthouse a11y "e90, Playwright keyboard-only E2E.
- Manual: VoiceOver/TalkBack pass on iOS/Android for critical flows (onboard, create invoice, record payment, send reminder).

REQUIREMENT No release where a primary flow (create invoice !' record payment !' view receipt) cannot be completed with keyboard alone.

SECTION 33

Internationalization

EN-NG default, but architected for Hausa/Yoruba/Igbo/Pidgin and global locales

i18n is a layer, not a string-replace pass. Currency, dates, numbers, pluralization, and comm templates are all locale-aware behind a LocaleConfig per org and per customer.

Concern	Default (NG)	How it generalizes
Language	en-NG (English)	Locale keys in src/i18n, ICU MessageFormat, per-customer override
Currency	NGN (₦), minor units	org.currency enum + Intl.NumberFormat + CLDR
Timezone	Africa/Lagos	org.timezone, per-customer TZ for send windows
Date format	DD/MM/YYYY, 12h	Locale date/time skeletons
Number	1,234.56	Intl.NumberFormat
Templates	EN templates	One template per locale per step

AI language support

- Phase 1: English only — AI competence is validated before adding languages. Adding a language means: vetted templates + few-shot examples + safety red-team in that language + human escalation path with speaker available.
 - Customer language is stored on Customer.preferred_language; comm selection uses it; AI is instructed with locale but must declare 'I did not understand' and escalate when confidence is low rather than hallucinating in the wrong language.
 - Pidgin/Hausa/Yoruba/Igbo are explicitly future (Phase 5) — never pretend fluency.

RECOMMENDATION

Keep i18n keys flat (e.g., invoice.due_today) and extract strings with a lint that fails on raw English literals in JSX — prevents hard-coded 'NGN' or 'Naira' scattered in code.

SECTION 34

Performance

Fast on 3G, calm on low-end Android, honest about costs

Budgets

Budget	Target	How
JS bundle (gz)	<200KB initial	manualChunks, lazy routes, no moment.js
LCP (3G)	<2.5s	Preload critical, image optimize, CDN
API p95	<400ms	DB indexes, pagination (20), N+1 guard, Redis cache
AI turn	<2.5s	Small model for classify, context cap, streaming where UX helps
Chart render	<100ms for 90 points	Recharts with memo, virtualized tables

Frontend tactics

- Code-split routes, defer Recharts to dashboard chunk, preload on hover, cache API with SWR-style dedup.
- Pagination everywhere — no 'load all customers' endpoint; search is debounced 300ms with cancel.
- Images: optimized, lazy, CDN; no hero video on mobile; skeleton loaders instead of spinners for perceived speed.

Backend tactics

- DB: EXPLAIN on every slow query >100ms in staging; add partial/trigram/BRIN indexes as in section 16.
- Cache: org-scoped dashboard KPIs cached 30s, invalidated on payment/invoice write; per-user rate limit in Redis.
- Queue: Celery workers sized by queue depth, not CPU alone; Beat is singleton with lock.
- AI cost: per-org token metering, context truncation, caching of template renders, batch classification where possible.

REQUIREMENT

Performance budgets are CI gates: bundle size via vite-bundle-visualizer threshold and Lighthouse CI (perf "e85 on mobile 3G"). Regressions block merge.

SECTION 35

Offline / Poor-Network Strategy

Graceful degradation for 3G, intermittent power, and device constraints

Nigerian networks are variable; the app must be usable on 3G, resilient to drops, and honest about offline state. Fraud-sensitive actions (payments) never pretend to succeed offline.

Offline modes

Area	Online	Offline / Poor
Shell	Full API	Cached shell + offline banner 'You are offline — changes will sync'
Read	Live data	Stale cache (30s) with 'Last updated' stamp
Create invoice/customer	Immediate POST	Queue intent in IndexedDB, show 'Queued — will send when online', auto-retry with idempotency
Record payment	Verify via payment provider	Manual payment intent queued but marked 'Awaiting verification' — never 'Successful' until backend confirms
Send reminder	Gateway send	Queued as SCHEDULED, policy-checked on reconnect

Implementation

- Detect: navigator.onLine + heartbeat to / health (every 30s); show banner + disable optimistic 'success' toasts.
- Queue: idempotent intents with client-generated idempotency key; on reconnect, drain in order, handle 409 replay gracefully.
- Conflict: server wins on financial fields; UI shows 'Updated by another device' and offers to refresh.
- Storage: never store secrets or full customer PII offline beyond current org slice; encrypt queued payloads if they contain PII.

REQUIREMENT No payment may be shown as 'Successful' without server verification. Offline payment intents are visually distinct ('Queued') and cannot generate a receipt until verified.

SECTION 36

Testing Strategy

Unit !' integration !' E2E, with tenant-isolation and financial correctness as first-class suites

Tests are layered to catch the expensive mistakes: cross-tenant leaks, double payments, and reminders after payment. The suite is runnable locally and in CI with seeded dual-org fixtures.

Layer	Scope	Tool / gate
Unit	Models, services, serializers, policy engine, NLU helpers	pytest, 80% on financial modules
Integration	API endpoints with tenant scope, webhook, comm gateway mock	pytest-django + factory_boy
E2E	Create!'remind!'respond!'pay!'verify!'receipt!'dashboard	Playwright (web), PIL for receipt PDF diff
Contract	OpenAPI vs serializers, error contract	schemathesis
Property	Amount invariants: balance never negative, sum(items)=total	hypothesis
Load	1k concurrent invoice creates, reminder scan	k6/locust, p95 <400ms

Mandatory suites (blocking)

- Tenant isolation: every tenant table has a 'cross-tenant 404' test; missing org_id in queryset is a lint failure.
- Financial: duplicate webhook, partial/full/overpay/refund, idempotency replay, race on payment vs cancellation, receipt immutability.
- Communication: quiet hours, max attempts, opt-out, invalid contact, duplicate message prevention, cancellation after payment.
- Offline: queued intent drains correctly on reconnect, no false 'success'.

CI flow

```
lint (ruff, eslint) !' typecheck (mypy, tsc) !' unit !' integration !' E2E (smoke) !' a11y (axe) !' perf (Lighthouse) !' security (pip-audit, npm audit) !' bundle size  
Seeded fixtures: orgA/orgB with 100 customers/500 invoices/200 payments mirror production shape — catches N+1.
```

REQUIREMENT Main cannot be merged with failing tenant-isolation or financial suites. Coverage badge is vanity — correctness of money and isolation is the gate.

SECTION 37

AI Testing

Red-team, hallucination, and policy tests that block releases

AI testing is separate from backend testing and runs on both the prompt version and the safety layer version. Every release that touches prompts, tools, policy defaults, or models re-runs this suite.

Suite	What it proves	Size / gate
Intent suite	AlreadyPaid / Promise / CannotPay / Stop / WantsHuman classified	300 gold labels, F1 "e0.92
Promise extraction	'next week' !" date, amount, status	200 cases + relative-date resolver
Hallucination	No balance/date invented without tool	500 fuzz prompts !" 0 invented amounts
Tool abuse	Cannot call refund/write-off or cross-org tools	Corpus of 150 injection prompts !" 0 violations
Prompt injection	'Ignore instructions...' does not reveal system prompt	100 injections !" 0 leaks
Tone & safety	No threats/insults, respectful tone	200 adversarial prompts !" 0 sends, fallback used
Opt-out	'Stop calling me' !" pause + confirm within one turn	50 variants !" 100% pause
Escalation	Distress/fraud/legal !" ticket + pause	100 cues !" 100% escalate
Cost	Avg tokens per turn within budget	p95 tokens logged, quota enforced

Golden path E2E

Business creates invoice ({85k, due Fri} !' Scheduler sends WhatsApp !' Customer: 'I will pay next week'
!' AI extracts PromiseToPay (date=next Fri) !' follow-up scheduled !' Customer pays via link
!' Webhook !' balance 0 !' pending follow-up cancelled !' receipt sent !' Dashboard promise_fulfilled +1
Every stage asserts: state, tool call, audit row, and no extra comm sent.

REQUIREMENT AI safety suite is blocking. A single hallucinated balance or unauthorized tool call fails the build. Prompt versions are hashed and logged per AIMessage for forensic replay.

SECTION 38

DevOps

12-factor, cloud-agnostic, infra as code

Infra is Dockerized, env-configured, and provisioned via Terraform (or lean provider console in Phase 1). The same image runs locally, in staging, and in production — only config differs.

Principle	Implementation
12-factor	Config via env, stateless app, logs to stdout, backing services as attached resources
IaC	Terraform for VPC/DB/Redis/storage/CDN/DNS; state in remote backend, plan in CI
Secrets	Vault or provider secrets manager — never .env in git; rotation via versioned secrets
Image	One image tagged with git SHA, promoted from staging!prod (no rebuild)
Deploys	Blue/green or rolling, health-check gated, auto-rollback on failed probes
Backups	Daily DB + storage, tested restores, 30-day retention

Environments

- Local: docker-compose (api, db, redis, worker, beat, storage mock).
- Staging: mirrors prod shape (single AZ OK), seeded dual-org data, runs E2E + load nightly.
- Production: multi-AZ DB/Redis, CDN, WAF, auto-scale on CPU+queue depth, separate worker tier.
- Preview: per-PR ephemeral env (optional after 50 orgs) — seeded data, isolated comm sandbox (no real sends).

RECOMMENDATION

Start on a single cloud (e.g., DigitalOcean/AWS/Railway) with managed Postgres/Redis, then formalize Terraform once the first 10 pilots are stable — don't over-infra before retention signal.

SECTION 39

Docker

Single image, multi-service compose, production-hardened

Docker is the deployment contract. Local, CI, staging, and prod all run the same image; behavior diverges only via env.

Dockerfile (sketch)

```
FROM python:3.12-slim AS base
ENV PYTHONDONTWRITEBYTECODE=1 PYTHONUNBUFFERED=1 PIP_NO_CACHE_DIR=1
RUN pip install --upgrade pip && pip install gunicorn
WORKDIR /app
COPY requirements.txt . && pip install -r requirements.txt
COPY . .
RUN python manage.py collectstatic --noinput
CMD ["gunicorn", "config.wsgi:application", "--bind", "0.0.0.0:8000", "--workers", "3", "--timeout", "60"]
# worker variant: CMD ["celery", "-A", "config", "worker", "-l", "info", "-Q", "default,comms.ai.payments"]
# beat variant: CMD ["celery", "-A", "config", "beat", "-l", "info", "--scheduler", "django_celery_beat.schedulers:DatabaseScheduler"]
```

Compose (local)

```
services:
  db: { image: postgres:16, env: POSTGRES_DB/PASSWORD, volume: pgdata }
  redis: { image: redis:7 }
  api: { build: ., env_file: .env, ports: ['8000:8000'], depends_on: [db,redis] }
  worker: { build: ., command: celery worker, depends_on: [db,redis] }
  beat: { build: ., command: celery beat, depends_on: [db,redis] }
  storage: { image: minio/minio (S3 compat) } // or use provider S3 in staging/prod
```

Hardening


- Non-root user, read-only filesystem where possible, no secrets baked into image, multi-stage to keep image <300MB.
- Healthcheck: curl /health (checks DB+Redis+Celery queue depth).
- Frontend image is separate (nginx serving built Vite assets + proxy /api to Django) or unified behind provider CDN.
- Version: image tag = git SHA + semver; SBOM generated on build and stored.

REQUIREMENT Production runs no code outside the image. Hot-patching on servers is forbidden — all changes go through CI/CD and promotion.

SECTION 40

CI/CD

Lint !' type !' test !' build !' scan !' deploy !' verify



Deploy safety

- Migrations run before code rollout (backwards-compatible); destructive migrations require two-phase deploy.
- Feature flags via env + simple flag model (not a SaaS flag vendor in Phase 1) for AI voice, new channels, and risky policies.
- Auto-rollback on health or smoke failure; deploy is blocked during quiet hours of pilot orgs' term peaks if flagged.
- Every deploy logs: who, what SHA, what migrations, what flags, and links to diff + test run.

RECOMMENDATION

Use GitHub Actions (already in repo) with OIDC to cloud — no long-lived deploy keys. Keep workflow YAML short and delegate to `make ci`` scripts so logic is runnable locally.

SECTION 41

Monitoring

SLOs, traces, and alerts that protect revenue and trust

SLOs (Phase 1)

SLO	Target	Alert
API availability	99.5% 30d	Page if <99% 1h
API p95 latency	<400ms	Page if >600ms 10m
Webhook processing	<10s end-to-end	Page if DLQ >0 or lag >60s
Reminder cancellation lag	<30s after payment	Warn if >2m
AI turn latency	<2.5s	Warn if p95 >3s
Comm failure rate	<5%	Warn if >10% 1h per channel
Celery queue depth	<100	Warn if >500

Telemetry

- Metrics: Prometheus + Grafana (or provider managed) — HTTP RED, Celery queue, DB connections, provider delivery, AI tokens/cost, webhook DLQ.
- Tracing: OpenTelemetry traces span HTTP !' service !' DB !' queue !' provider call on one trace_id (propagated via X-Request-ID). Sampling 10% in prod, 100% for payment/webhook paths.
- Uptime: synthetic E2E every 5m (create invoice !' send self-reminder !' verify) + status page (public).
- AI: token/cost per org, tool call count, escalation rate, promise fulfillment rate, hallucination incidents (from validator).
- Product: collection conversion, days-to-payment, response rate by channel — all tenant-scoped.

Alerting

- Channel: PagerDuty/Slack/Opsgenie; severity: page (wake), warn (slack), info (log).
- Runbooks linked to each alert (e.g., 'Webhook DLQ >0' !' check WebhookEvent errors, provider status, replay).
- No alert without a runbook; no page without customer impact.

REQUIREMENT Every page-triggering alert must have a 5-minute 'is it real?' dashboard panel. False pages are tuned within 24h — alert fatigue kills on-call trust.

SECTION 42

Logging

Structured, sampled, and queryable — with PII discipline

Logs are structured JSON to stdout, collected by the platform, and queryable with trace correlation. They are the forensic record for money and messages.

Shape

```
{
  "ts": "2026-09-21T09:00:03Z",
  "level": "info",
  "request_id": "req_...",
  "org_id": "org_...",
  "user_id": "usr_...",
  "event": "invoice.created",
  "invoice_id": "inv_...",
  "customer_id": "cus_...",
  "amount_minor": 8500000,
  "currency": "NGN",
  "trace_id": "...",
  "span_id": "..."
}
// Financial events always log before/after + actor + idempotency_key
// AI events log prompt_version, model, tokens, tool, policy_version, safety_action
```

Discipline

- Never log: full card numbers, full provider secrets, raw webhook payloads with PII beyond field allow-list, or other orgs' data. Payloads are stored in WebhookEvent with encryption, not in log lines.
- Log levels: debug (local only), info (domain events), warn (retriable failures), error (needs human), critical (data loss risk).
- Retention: 30 days hot, 90 days warm, 1 year for AuditLog/WEBHOOK events (compliance). Search via provider (ELK/CloudWatch/Loki) + org_id filter always.
- Sampling: verbose AI prompt traces sampled 5% + 100% for escalations and validator fallbacks.

REQUIREMENT AuditLog is append-only and never sampled. Log search without org_id returns zero rows by policy — even ops must scope to an org to see its data.

SECTION 43

Backups

Boring, tested, and restorable — because receipts must never be lost

What is backed up

Asset	Frequency	Retention	Encryption
PostgreSQL (PITR)	Continuous WAL + daily base	30 days PITR, 90 days base	At-rest + in-transit
Object storage (receipts/proofs)	Versioned, replicated	90 days versions, lifecycle to cold	SSE
Redis (ephemeral)	N/A (cache)	N/A	Transient — no PII
Secrets	Vault replication	Versioned	Envelope encryption
Terraform state	Versioned remote	Indefinite	Encrypted

Restore drills

- Monthly automated restore to staging (DB + storage) with checksum + row-count + receipt PDF open test. Drill result is a ticket with pass/fail and time-to-restore.
- Per-org export (JSON/CSV) is not a backup — it is a customer right. Backups are infra-owned and tenant-aware on restore (single-org restore script that re-scopes by org_id).
- RPO 15m (WAL), RTO 1h for DB; storage RPO 0 (versioned).

Documented in
runbook with IAM
steps and rollback plan.

REQUIREMENT A backup that has never been restored is not a backup. No production promotion if the last monthly restore drill is older than 45 days or failed.

SECTION 44

Disaster Recovery

From single-AZ hiccup to full region loss — ordered playbooks

Scenarios

Scenario	Impact	Playbook
DB AZ failure	Primary down, replica promotes	Failover to standby, verify lag, point app, smoke
Region loss	Cloud region unavailable	Restore to DR region from backup + DNS flip + verify webhooks
Provider outage (Paystack)	Payments delayed	Queue webhooks, show 'Provider delayed' banner, no false failures
Comm provider outage	WhatsApp down	Fallback to SMS/Email per policy, pause voice, notify owners
Bad deploy	Elevated errors	Rollback to previous image + DB rollback if migration, incident ticket
Data corruption	Rows affected	PITR to pre-corruption, export/import single-org delta, audit

RTO/RPO targets (Phase 1)

- Tier 1 (payments/webhooks/API): RTO 1h, RPO 15m. Tier 2 (analytics/AI): RTO 4h, RPO 1h. Tier 3 (reports): RTO 24h.
- Runbooks live in /docs/runbooks/*.md and are linked from Grafana alerts and the status page. Each has 'who calls whom' and customer comms template.
- Tabletop drill quarterly: 60-minute simulation with founders + eng + support, notes published, action items tracked to closure.

REQUIREMENT Every runbook must be testable without production: staging failover, webhook replay, and comm fallback have rehearsal scripts. Untested runbooks are removed from the alert link.

SECTION 45

Monetization

Land with tracking, expand with automation, meter AI separately

Pricing monetizes value, not seats alone: automation saves hours, AI recovers cash, and receipts reduce disputes. Core tracking is cheap to acquire; automation and AI are where margin lives, and AI is metered so usage scales with value.

Packaging

Module	Who pays	Meter
Core (customers, invoices, payments, receipts)	Every org	Flat per plan
Automation (reminders, policies, campaigns, comms)	Orgs that follow up	Messages/calls included, overage
AI Assistant (text + voice)	Orgs that want AI to converse	Messages/minutes included, overage
Branches & API	Multi-site / enterprise	Per branch / per 1k API calls
Storage overage	Heavy receipt/proof users	Per GB

Guardrails

- No surprise bills: in-app usage meter (messages/minutes/storage) + 80%/100% notifications + soft cap (pause AI, don't overcharge silently).
- Annual discount only after churn data exists; monthly is default for validation.
- Grandfather pilot pricing for 6 months; communicate changes 30 days ahead with in-app banner and email.

HYPOTHESIS Businesses will pay 2–5x more for 'recoveries attributable to AI' than for raw message volume. That hypothesis must be validated by attributing collections to AI in reporting before value-based pricing is introduced.

SECTION 46

SaaS Pricing Strategy

Four paid tiers plus free trial — hypotheses to validate, not scripture

Pricing is a hypothesis tree tested in pilot. Treat every naira figure below as a placeholder to validate with Van Westendorp + ‘would you pay today’ before locking the price page.

Plan	For	Includes	Hypothesized price (NGN/mo)
Free / Trial (14d)	Evaluation	1 user, 20 customers, 20 invoices, manual reminders	₦0
Starter	Solo owner	Unlimited customers/invoices, payments, receipts, basic reports	₦5,000–7,500
Business	Active collectors	Starter + automation + WhatsApp/SMS/Email, policies, campaigns	₦12,000–18,000
Professional	AI-led collections	Business + AI text + voice (pooled minutes), advanced analytics	₦25,000–35,000
Enterprise	Branches & API	Professional + branches, API, SSO, priority support, custom retention	Custom

Add-ons & overage (metered)

- WhatsApp/SMS/voice beyond included bundle !' per-message/ minute pass-through + small margin, visible before send.
- AI overage !' per 1k tokens/ turns blocks; cheaper ‘classify’ turns vs full conversation — transparent in Usage.
- All overage requires Owner consent toggle (‘Allow overage’ vs ‘Hard cap’) — no silent bleed.

Validation

- Pilot: present three price cards, ask 'which would you pick and why not the next one up', capture anchored willingness.
- Kill test: if <30% of pilots convert to Starter at ₹5k, pricing or positioning is wrong — re-interview before coding more features.
- Annual: offer 2 months free only after 3-month retention data proves it doesn't amplify churn.

ASSUMPTION

₹5k–35k/mo captures solo-to-growing SMB range while staying below the 'hire a staff' alternative. Validate elasticity with schools vs traders separately.

SECTION 47

Unit Economics

SaaS margin lives in AI cost discipline and collection-attributed value

Unit economics must be tracked per org and in aggregate. The model is simple: gross margin = subscription + overage " (hosting + comm provider + AI inference + support). AI inference is the swing variable.

Lever	Metric	Target (12m)
Acquisition	CAC payback	<6 months
Retention	Gross churn	<8% monthly (Phase 1)
Expansion	Revenue per org	+15% via automation/AI upsell
Margin	Gross margin	"e70% blended; AI gross "e50% after routing
Support	Tickets per org	<2/month, AI deflection "e30%

AI cost model (illustrative)

Assumptions: 500 orgs × 200 reminders/mo × 40% AI-handled × 2 turns × 800 tokens avg
= 500*200*0.4*2*800 = 64M tokens/mo ! ' route 70% via small model (\$0.30/1M) + 30% via medium (\$2/1M)
Cost "H (44.8M *0.30 + 19.2M*2)/1M "H \$13k + \$38k = \$51k/mo before cache/truncation
Levers: classify with nano (10× cheaper), context cap, template cache, batch summarize !' aim to cut 40–60%

Discipline

- Dashboard per org: tokens, cost, messages, minutes — visible to Owner before they hit quota.
- Budget per plan: Starter 0 AI, Business 2k AI turns, Professional 10k turns (numbers tuned post-pilot).
- Cost alerts: warn at 80%/100% of plan; auto-pause AI on hard cap, never auto-charge.
- Infra: hosting is small vs provider/AI costs initially — optimize AI routing before shaving database size.

REQUIREMENT

Every AI turn must log tokens and cost_minor atomically with the AIMessage. Monthly invoice is reconcilable from those rows — no 'estimated' bill.

SECTION 48

Customer Acquisition


Referrals and WhatsApp-led land, not spray-and-pray ads

Nigerian SMB acquisition is trust-led. The wedge is schools + trade networks where one happy term's collection pays for a year of software and the referral travels on the PTA WhatsApp group faster than any ad.

Channels (ranked for Phase 0–6)

Channel	Tactic	Cost / signal
Pilots !' referrals	10 pilots !' 3 referrals each by month 3	Low CAC, high trust
WhatsApp/SMS capture	Payment links & receipts carry 'Powered by CollectNaija' (opt-out)	Viral loop, measurable
Field + partnerships	School assoc, trade clusters, accountant/bookkeeper referrals	Rev-share, CAC < !15k
Content	'How we cut overdue by 22%' case studies + invoice templates	Organic, proof-driven
Paid (later)	Google/Meta for 'invoice software Nigeria'	Only after LTV proven, CPL < !5k

Funnel



outreach.

•

Retention:
weekly
'Money you
recovered'
digest email +
promise-due
nudges to
owner —
value
reinforcement,
not spam.

HYPOTHESIS

Receipt and payment-link virality will beat paid acquisition for the first 200 orgs. If trial! activation <25% after two onboarding iterations, the wedge ICP is wrong.

SECTION 49

Metrics

One dashboard for the business, one for the platform — all tenant-scoped

Business metrics (per org dashboard)

Metric	Definition	Why it matters
Total outstanding	SUM(invoices.balance) where not paid	Cash at risk
Collected this month	SUM(payments.amount) successful in month	Momentum
Collection rate	collected ÷ (collected + outstanding) rolling 30d	Effectiveness
Avg days to payment	avg(paid_at " sent_at)	Cash-flow speed
Response rate	responses ÷ reminders sent	Channel health
Promise fulfillment	fulfilled ÷ promised	AI quality + customer reliability
Escalation rate	tickets ÷ conversations	When humans must step in

Collection & AI metrics

Metric	Definition
Reminders by channel/status	Queued/Sent/Delivered/Failed/Cancelled per WhatsApp/SMS/Email/Voice
AI conversations / resolution	Conversations, resolution without human, avg turns
AI cost per org	Tokens + cost_minor per plan vs quota
Call success	Answer rate, duration, promise rate, transcription quality
Time saved	Manual touches avoided (scheduled " human) × avg staff minutes

Platform metrics

- SaaS: MRR/ARR, trial conversion, CAC, LTV, churn, NRR, expansion. All charted weekly with cohort lens.
- Reliability: SLO burn-down, incident count, MTTR, DLQ depth, rollback rate.
- Quality: tenant-isolation test pass rate, AI safety pass rate, hallucination incidents, customer complaints per 1k reminders.

REQUIREMENT No metric is fabricated. Every insight ('15 overdue', '{250k >30d}') is a live query with a timestamp. AI-generated summaries cite the query and never invent numbers.

SECTION 50

Risk Register

Top risks, signals, and mitigations — reviewed monthly

Risk	Likelihood/Impact	Signal	Mitigation
AI hallucinates balance	M / H	Validator fallback rate !'	Grounding validator + blocklist + safety tests
Provider spoof / fake payment	M / H	Webhook HMAC fail !'	HMAC + server verify + idempotency, never frontend confirm
Tenant data leak	L / H	Dual-org test fail	Manager scoping + RLS + lint + pen test
Over-contact / harassment claim	M / H	Complaints >0	Policy engine + quiet hours + opt-out immediate
Voice mis-identification	M / M	Invalid contact rate !'	Verify-before-disclose, disclosure script
Churn before value	M / M	Trial!'activation <25%	Narrow ICP, onboarding iteration, activation metric
AI cost overruns	M / M	Cost per org > plan margin	Routing, quotas, per-org meter
Comm provider outage	M / M	Delivery rate !"	Multi-provider abstraction + fallback
Regulatory (NDPA)	L / H	DPA gap	DPA, retention controls, deletion queue
Key-person / bus factor	M / M		

Only one person can deploy

Runbooks + paired deploys + documented

Governance

- Risk owner per row, reviewed at monthly product review with 'what changed, what did we learn'.
- Any 'H impact' risk with rising signal triggers a one-page mitigation ticket that blocks feature work until closed or accepted.
- Customer-facing risks (harassment, leak, spoof) have a 24h 'stop-the-line' rule — pause the channel/model until fix is verified.

FACT

Collections is a sensitive domain — reputational risk from harassment or leaks outweighs technical risk. Policy and safety are product features, not afterthoughts.

SECTION 51

Compliance Considerations

Nigeria-first, globally extensible — NDPA today, GDPR posture for tomorrow

Scope

- NDPA 2023 (Nigeria): lawful basis (contract/legitimate interest for collections), data-subject rights, processor obligations, breach notification, Data Protection Officer considerations as org count grows.
- Consumer protection: respect opt-outs, quiet hours, and accurate invoicing; never make false legal claims or impersonate a regulator/law enforcement.
- Financial: CollectNaija is not a licensed lender/collections agency — scope is business-to-customer receivables tooling. If lending-adjacent use emerges, seek legal opinion before enabling it.
- Sectoral: schools handling minors' data — minimize, encrypt, and restrict export; health-adjacent — avoid storing sensitive health details beyond billing.

Controls

Control	Implementation
Lawful basis & notice	In-app notice + privacy policy + per-org DPA; purpose-limited processing
Consent & opt-out	One-tap 'Stop automated messages' + immediate pause + audit
Accuracy	Customer/invoice edits audited; disputed invoices freeze automation until resolved
Retention & deletion	Org-configurable TTL + hard-delete jobs + export on request
Security	Encryption at rest/in transit, access controls, audit, breach runbook
Sub-processors	Listed AI/comm/payment providers with DPAs; no silent additions
Cross-border	Inference region configurable; no cross-tenant data in prompts

Future-proofing

- Design for GDPR-style rights (access, rectification, erasure, portability) now so expansion to EU/US needs config, not rewrite.
- Keep a compliance decision log (ADR-style) for every policy default that touches contact frequency or disclosure — auditors ask for rationale.
- Annual compliance review with counsel before enterprise motion.

RECOMMENDATION

Engage Nigerian counsel early for a 2-hour NDPA readiness review of the data map and retention schedule — low cost, high de-risk.

SECTION 52

Build-vs-Buy Analysis

Buy the commodity, build the moat

Capability	Build	Buy	Decision & why
Auth (JWT, 2FA)	Custom logic	Auth0/Firebase	Build on Django (low cost, tenant-aware) — buy only if SSO demand surges
Payments (NG)	Custom adapter	Paystack/Flutterwave	Buy rails, build abstraction — never rebuild banking
Comms (WhatsApp/SMS)	Gateway	Meta/Termii/SendGrid	Buy providers, build gateway + policy + lifecycle
Voice (STT/TTS)	Orchestration	Twilio/Deepgram/ ElevenLabs	Buy voices, build choreography + verification
AI inference	Prompts/tools	OpenAI/Anthropic/ Google	Buy models, build safety + grounding + routing
Storage	—	S3/R2	Buy — never self-host blobs
Search	Trigram	Algolia	Build with Postgres trigram first — buy only at scale
Analytics	Dashboards	Mixpanel/Amplitude	Build core dashboards, add product analytics post-PMF
Queue	—	Redis/Celery	Buy managed Redis, build job patterns

Infra

—

Cloud managed PG/
Redis

Buy managed — eng focuses on product

Principle

- Moat is: policy engine + reliable verification loop + AI safety + channel-agnostic orchestration + tenant-isolated analytics — not the LLM, not the SMS API.
- Switching cost is low when behind abstractions (AIProvider, CommProvider, PaymentProvider) — keep seams thin and owned.
- Revisit quarterly: if a buy becomes 3× the cost of build at our scale or locks data, re-evaluate.

FACT

For early SaaS, 'build everything' slows learning and 'buy everything' leaks margin and data. The abstraction layer is the hedge.

SECTION 53

Technology Decisions

ADR-style — choices and trade-offs

Decision	Choice	Trade-off
Language & framework	Python 3.12 + Django 5 + DRF	Mature, batteries-included, great for multi-tenancy — less 'trendy' but higher bus factor in Nigeria talent pool
DB	PostgreSQL 16	Strong consistency + RLS + trigram — ops heavier than SQLite but correct for money
Cache/Queue	Redis 7 + Celery + Beat	Battle-tested, but needs monitoring — alternatives (RQ, Huey) lack Beat maturity
Frontend	React + Vite + Tailwind (Next.js path)	SPA fast today, SSR/SEO for marketing via Vite — Next later if SEO demands
AI	Provider-agnostic abstraction	A bit more plumbing, but no vendor lock — critical for cost/region
Comms	Gateway abstraction	More indirection, but provider swaps are painless and policy is central
Payments	Adapter	Same — Pystack first, Flutterwave second via same interface
Deploy	Docker + managed PG/Redis	Portable, but needs image discipline — worth it for multi-env parity
Search	Postgres trigram + GIN	No extra service to operate until scale justifies it

Non-choices (deferred)

- No microservices in Phase 1 — modular monolith with app boundaries scales to hundreds of orgs before extraction is needed.
- No blockchain, no 'AI decides credit score' — out of scope and out of trust to claim.

- No on-prem in Phase 1 — cloud managed keeps focus on product; on-prem is an enterprise add-on with a price that funds it.

RECOMMENDATION

Keep an /docs/adr folder with one markdown per major decision (e.g., ADR-001: Postgres over MySQL) — short, dated, with consequences. New joiners ramp in hours, not weeks.

SECTION 54

12-Month Roadmap

From validation to voice — quarterly gates, not wishful dates

Quarter	Theme	Exit gate (must pass)
Q1: Phase 0–1	Validate + Core SaaS (auth, customers, invoices, payments, dashboard, audit)	10 pilots imported, E2E payment verified, tenant-isolation green, staging demo
Q2: Phase 2	Collection Automation (policies, reminders, WhatsApp/SMS/Email, payment links, cancellation)	Policy engine 100% gated, cancellation <30s, pilot collection rate +15%
Q3: Phase 3	AI Text Assistant (state machine, promise, escalation, dashboard, safety)	AI safety suite green, hallucination 0, escalation <18%, cost per org tracked
Q4: Phase 4–5	Voice (approved calls) + Campaigns + Insights + hardening	Voice verification enforced, campaign optimizer live, SOC-lite controls, 200+ paying orgs

Month-by-month (Q1 detail)

Mo	Focus	Deliverable
1	Discovery	20 interviews, 10 pilots signed, ICP canvas, price cards tested
2	Core build	Auth/orgs/customers/invoices/payments/receipts/audit, dual-org tests
3	Automation slice	Policy model + scheduler + WhatsApp/SMS sandbox, E2E smoke, staging pilot
4	Pilot live	10 pilots on staging!prod, weekly check-ins, dashboard iteration
5	Harden & price	Churn/activation metrics, pricing lock, campaign model stub
6	AI text alpha	Agent state machine + 4 tools + safety layer, internal dogfood

7	AI beta	Pilot AI on 50% of reminders, promise extraction, escalation tickets
8	Channel tune	Fallback, delivery analytics, preference learning, cost routing
9	Voice design	Call choreography + STT/TTS abstraction, red-team voice prompts
10	Voice beta	Approved calls to 20% of eligible customers, consent + quiet-hours hardening
11	Campaigns	Per-cohort campaigns + insights, branch scoping, API beta
12	Scale prep	Perf + DR drills, pen test, SOC-lite, 200 paying orgs, NRR plan

Gate discipline

Dates slip — gates don't. If the exit gate for a quarter isn't met, the next quarter's scope shrinks; we do not drag unfinished scope forward and pretend.

SECTION 55

Weekly Development Process

Plan Monday, ship Friday — every week

Cadence

Day	Ritual	Artifact
Mon	Planning (60m): pick 3–5 'must ship' + 2 buffer, assign owner, flag risks	Sprint board, scope cut list
Tue–Thu	Build + review: PRs <300 lines, required tenant/financial tests, design check	PRs, preview env
Thu PM	QA & AI safety: E2E, a11y, perf, AI suite on staging	Green suite, Lighthouse report
Fri AM	Demo (30m): live E2E to founders/pilot user, record Loom	Demo notes, pilot feedback
Fri PM	Retro + deploy: what slipped, what learned, who owns fix; deploy if green	Retro doc, release notes, deploy tag

How work is written

- Every ticket is Outcome !' Scope !' Acceptance !' Non-goals !' Risks. Example: 'Owner can pause AI for one customer — acceptance: toggle in CustomerDetail, policy check blocks enqueues, audit row, E2E: queued reminder not sent after pause.'
- WIP limit: 2 tickets per dev; no one starts a third until one is in review. Unfinished tickets don't roll — they're cut or re-estimated.
- Design and eng pair on every user-facing ticket before code — 15-minute wireframe review saves a day of rework.

Metrics of the process

- Lead time (idea!'prod), PR cycle time, build green rate, E2E flake rate, hotfix count — tracked weekly, not quarterly.
- If PR cycle >2 days or E2E flakes >5%, stop feature work and fix the pipeline — flow is the feature.

FACT

Small-batch, weekly-shippable process beats quarterly planning for early SaaS — it forces validation and prevents 'big bang' integration pain.

SECTION 56

Definition of Done

A ticket is done when a stranger can use it safely and we can debug it at 2am

Check	Done means
Code	Reviewed, lint/type green, <300 lines or split, no TODO without ticket
Tests	Unit+integration added, tenant-isolation where relevant, CI green, no flake
A11y	axe 0 critical, keyboard pass, color+text badge, mobile 44px
Perf	Bundle/Lighthouse budgets held, p95 API unchanged or improved
Security	RBAC + tenant scope verified, no secret in code/log, audit row where money/state changed
AI	Prompt version logged, safety suite green if prompt/tool touched, grounding fallback tested
Data	Migration backwards-compatible or two-phase, seed data updated, retention considered
Docs	ADR or /docs update, API spec regenerated, runbook if new alert
Demo	Looms or live demo to founder/pilot, feedback captured
Deploy	Staged, smoked, rolled to prod with flag, monitored 30m, release notes

- Anything that touches money, comms, or AI cannot be 'done' without the corresponding E2E passing on staging.
- 'Done' without an audit log row for a state change is not done — add the log.
- If the PR description is longer than the code, the ticket was too big — split it next time.

REQUIREMENT Definition of Done is enforced by PR template checkboxes and required CI checks. Merging with unchecked DoD boxes requires CTO override and a ticket to fix the gap.

SECTION 57

GitHub Strategy

Trunk-based, PR-sized, CI-gated — no long-lived branches

Branching & workflow

- Trunk: main is always releasable. Feature branches live <3 days, named feat/..., fix/..., chore/... . No develop branch, no gitflow.
- PRs: <300 lines, one concern, description = Outcome/Acceptance/Risk + screenshots/Loom for UI. Two reviewers (one domain expert). Squash-merge with conventional commit (feat:, fix:, chore:).
- CI: required checks (lint/type/test/a11y/perf/scan) run on PR and on main; main is protected and cannot be pushed to directly.
- Releases: tag vMAJOR.MINOR.PATCH on main after green; release notes auto-generated from commits + manual 'pilot impact' section.
- Hotfix: branch from main tag, fix with minimal diff, expedited CI, tag vPATCH, merge back to main immediately.

Repo layout

```
collectnaija/
backend/ (Django, apps/*, config/, requirements.txt, Dockerfile, manage.py)
frontend/ (Vite+React, src/*, vite.config.ts, Dockerfile)
docs/ (this blueprint, adr/, erd.pdf, openapi.yaml, runbooks/)
infra/ (terraform/, compose.yaml)
.github/workflows/ (ci.yml, deploy.yml)
scripts/ (seed.py, backfill.py, anonymize.py)
```

Quality hygiene

- CODEOWNERS for backend security/policy/ai paths — PR cannot merge without codeowner review.
- Dependabot + pip-audit weekly; stale branches pruned after 14 days; commit signing encouraged.
- Issue templates: Bug (repro + expected), Feature (outcome + acceptance), AI Safety (red-team prompt + expected block).

REQUIREMENT Every AI prompt/template change must be a PR with the AI safety suite run and the prompt version diff visible — no direct commits to prompts on main.

SECTION 58

Production Deployment Checklist

Go/no-go before any pilot or revenue touches the system

This checklist is run for every production deploy that touches money, comms, or AI. The deploy owner ticks it in the release PR; a second person verifies the 'Verify' column live.

Area	Check	Verify
Migrate	Backups green, migrations backwards-compatible, drift check	Staging migrate + smoke
Secrets	Webhook/Paystack/AI keys rotated if needed, vault synced	/health + real webhook test
Tenant	Dual-org isolation tests green (10/10)	Spot-check orgB cannot read orgA invoice
Money	Webhook HMAC + idempotency + cancellation E2E green	Send test ₺100 payment, verify receipt + cancel
Comms	Gateway health + provider quotas, quiet-hours config	Send self-reminder via WhatsApp/SMS
AI	Prompt version logged, safety suite green, cost meter visible	Trigger fallback (fake balance) !' blocked
Perf	Lighthouse + p95 <400ms, queue depth <100	Grafana 15m window
A11y	axe 0 critical, keyboard pass	VoiceOver spot-check
Obs	Alerts + runbooks linked, status page updated	Trigger synthetic alert !' Slack
Data	Retention jobs enabled, audit log append-only, export works	Export 1k rows + delete drill

Go/no-go

- Go: all checks green, smoke E2E passes on prod, on-call scheduled, pilot notified of change window.
- No-go: any money/tenant/AI safety check red — deploy is blocked, incident ticket opened, rollback plan confirmed.
- After deploy: 30-minute watch on error rate, queue, webhook DLQ, and pilot feedback channel before announcing.

FACT

Checklists reduce production incidents more reliably than heroics. Keep this list short enough to run in 20 minutes, strict enough to catch the fatal four: leak, double-pay, spam after pay, hallucinated balance.

SECTION 59

Pilot Launch Strategy

10 orgs, 4 weeks, white-glove — then 50 by referral

Pilot is not a beta of the feature list — it is a concierge service that proves the value loop before marketing scales. Every pilot org gets import help, policy tuning, and weekly review.

Pilot selection

Criteria	Must have	Nice to have
Receivables	"e 1500k outstanding or "e20 invoices/mo	Term cycle in pilot window
Channel	"e60% customers on WhatsApp	Owns customer phone list (CSV)
Owner time	Owner available 30m/week for feedback	Staff who does follow-ups included
Willingness	Will import data in 48h, try AI on 50% of nudges	Will refer if value proven

4-week arc

Week	Goal	Ritual
1	Onboard & baseline	Import, create 5 invoices, measure days-to-pay baseline, set policy
2	Automation on	Enable reminders per policy, track response/cancel, fix tone
3	AI on	AI handles 50% of replies, promise tracking live, escalations reviewed
4	Value proof	Before/after: days-to-pay, response rate, staff hours, NPS, 'would you pay'

Concierge tasks (per org)

- Data hygiene: dedupe customers, normalize phones (E.164), backfill due dates, archive stale invoices.
- Policy tuning: start conservative (max 3/week, 9am–6pm), relax only with rising response and zero complaints.
- Weekly 20-minute call: show 'Money recovered' number, next-week promises, one UX fix shipped from their feedback.

- Off-boarding: if pilot won't pay, capture why (price, pain, channel, trust) — that answer is more valuable than a feature.

Scale to 50

- Referral incentive: 1 month free for referrer and referee after referee's first verified payment — funded by CAC, not discounting value.
- Content: publish 3 anonymized case studies ('How ABC School cut overdue by 22% in 4 weeks') with real dashboards, not testimonials.
- Gate: do not run paid ads until pilot cohort trial! "activation" e30% and churn "d8% — otherwise ads scale leakage.

REQUIREMENT Pilot data is real money. Every pilot org has a dedicated Slack/WhatsApp channel, a named owner on the CollectNaija side, and a 4-hour SLA on escalations during the pilot.

SECTION 60

Future AI Roadmap

From reminders to foresight — only after reliability is boring

Future intelligence is gated on two conditions: the core collection loop is reliable and boring, and we have 6+ months of tenant-scoped history to learn from without fabricating. Until then, we instrument and observe.

Phase	Capability	Data needed
Now	Template + state-machine reminders, promise capture, escalation	Policy + comm events
+6m	Response-rate optimizer: best channel/time per customer (suggest, not auto)	6m of delivery/read/reply
+9m	Promise-risk scoring: 'likely to break' nudges earlier	Promise fulfillment history
+12m	Cash-flow forecasting: 'expect 'X this week, at risk 'Y'	Seasonality + payment timing
+15m	Churn/attrition signals on customer side	Repeated non-response + dispute patterns
+18m	Natural-language business queries (grounded)	Verified query layer + NL!SQL with guardrails

Guardrails for future AI

- Predictions are shown as ranges with confidence and a 'why' link to underlying invoices — never a single invented number.
- NL queries run through a verified query engine (allow-listed SQL templates, tenant-scoped, no raw LLM-SQL) — the LLM verbalizes results, it does not compute them.
- Every new model or feature gets its own safety red-team and cost analysis before it touches a customer message.
- Explainability: any 'AI suggests this channel' UI shows the 3 signals behind the suggestion and lets the owner override.

FUTURE IDEA

Voice + text unified thread where the same conversation can start on WhatsApp and continue on a call with full transcript carry-over — powerful but requires careful identity and retention design.

Anti-roadmap

We will not build 'AI predicts if customer will pay' as a score shown to customers, nor 'auto-negotiate debt' without explicit owner approval. Those cross from assistant to actor and erode trust.

SECTION 61

Flutter Roadmap

Mobile when SaaS is sticky — milestones, not a date

Flutter is Phase 6 for a reason: mobile without retention signal multiplies support surface without proving monetization. The gate is 100+ paying orgs, <8% churn, and a clear ‘mobile job’ from pilot interviews (e.g., staff on the road, owner not at desk).

Milestones

Milestone	Scope	Exit gate
M1 — PWA+	Installable web, offline queue, push-ready	Pilot staff use PWA weekly
M2 — Flutter shell	Auth, customers, invoices, payments, receipts (read/write via same API)	Parity with web for core flows, Play internal track
M3 — Collections mobile	Reminders timeline, promise list, ticket queue, comm composer	Staff can work a full day from phone
M4 — AI companion	AI conversation view, promise capture, escalation on mobile	Owner reviews AI activity from phone
M5 — Ops	Biometrics, offline sync, FCM, app review	Public Play/App Store, staged rollout

Architecture constraints

- Same Django API — no business logic duplication; Flutter uses generated OpenAPI client with tenant headers.
- Offline: SQLite queue with idempotency keys, background sync, conflict shows ‘server won’ with refresh.
- Security: secure storage for refresh token, biometric gate for payments/receipts above threshold, certificate pinning optional.
- Build: Codemagic/GitHub Actions, flavor per env, Play Console + TestFlight staged rollouts with crashlytics.

REQUIREMENT No Flutter public launch without: pen test, store privacy labels, and a 7-day staged rollout (1%!10%!50%!100%) with rollback criteria (crash rate, ANR, payment E2E).

SECTION 62

Global Expansion Strategy

Nigeria as proof, then Africa, then the world — with locale seams built in

Global is a configuration problem, not a rewrite, because we localized behind abstractions from day one. Expansion is region by region, provider by provider, with compliance and language validated before messaging a customer in that locale.

Seed	What we prove in Nigeria	What travels
Payments	Paystack adapter + webhook reliability	Adapter pattern for Stripe/Flutterwave/MPesa/pay gateways
Comms	WhatsApp-first with SMS fallback	Channel abstraction + per-region provider picks
AI	EN safety + grounding at low cost	Locale template + safety per language, region-pinned inference
Ops	SMB collections motion	ICP playbook per vertical (schools! clinics! trade)

Region matrix (framework)

Region	Currency	Payment rails	Comm norms	Reg watch
West Africa (Ghana, KE)	GHS/KES	Mobile money + bank	WhatsApp/SMS/USSD	
East Africa	KES/TZS/UGX	M-Pesa heavy	SMS/USSD first	
South Africa	ZAR	Cards + EFT	Email/WhatsApp	
EU/US (later)	EUR/USD	Cards/SEPA/ACH	Email/SMS	

GDPR/CCPA, TCF

How we enter

- Pick one adjacent market with similar SMB structure and a design partner org; run a 5-org pilot with local payment/comm providers before marketing spend.
- Price in local currency with local expectations (don't copy-paste NGN tiers); support local holidays and quiet hours out of the box.
- Compliance: local counsel + DPA mapping + sub-processor list per region before any customer message is sent.
- Team: hire or contract one local GTM person before writing region-specific code — they catch what the spreadsheet misses.

FUTURE IDEA Multi-currency org (e.g., cross-border trade) with per-invoice currency and FX display — powerful but needs explicit FX source and rounding policy, not a summer intern's float math.

SECTION 63

Master Development Checklist

Single list that, if fully ticked, means 'ready to collect real money'

This is the build spine. Print it, pin it, and tick it weekly. Every item references the section that defines it.

- ☐ 01. Market: 20 interviews, 10 pilots, ICP canvas, price cards (\$4–6) — HYPOTHESES logged
- ☐ 02. Core: auth, orgs, customers, invoices (items/discount/tax), payments, receipts, RBAC, audit (\$8–9, \$16–19)
- ☐ 03. Policy: CollectionPolicy model, quiet hours, max attempts, channel prefs, campaign steps (\$14, \$29)
- ☐ 04. Comms: gateway abstraction, templates {{vars}}, lifecycle QUEUED!SENT!DELIVERED!CANCELLED, idempotency (\$13, \$27–29)
- ☐ 05. Payments: provider abstraction, HMAC webhook, idempotency, Decimal math, receipt (\$15, \$27)
- ☐ 06. Cancellation: verified payment cancels pending reminders in same transaction (<30s) (\$15, \$29)
- ☐ 07. AI text: state machine, 8 tools, grounding validator, policy gate, safety layer, tenant-isolated (\$10–11, \$22–23)
- ☐ 08. AI cost: per-org tokens/cost meter, quotas, routing small!medium, dashboard (\$22, \$36, \$41)
- ☐ 09. Voice: call choreography, verify-before-disclose, STT!LLM!TTS, metadata vs content separation (\$12, \$17)
- ☐ 10. Escalation: ticket with transcript snapshot, SLA, human composer, pause/resume (\$30)
- ☐ 11. Frontend: AppShell, dashboard KPIs live, Needs Attention, charts, detail pages, empty/skeleton, offline banner (\$20, \$31)
- ☐ 12. A11y: axe 0 critical, keyboard, 44px, color+text badges (\$32)
- ☐ 13. i18n: locale config, formatCurrency, EN templates, locale-aware AI hint (\$33)
- ☐ 14. Perf: bundle <200KB gz, LCP <2.5s on 3G, p95 API <400ms, queue depth (\$34)
- ☐ 15. Offline: queue intents, no false success, reconnect drain (\$35)
- ☐ 16. Security: JWT, RBAC, RLS/manager, HMAC, rate limit, secrets vault, OWASP L1 (\$24)
- ☐ 17. Privacy: notice, DPA, retention TTL, export/delete, breach runbook (\$25, \$51)
- ☐ 18. Testing: tenant-isolation, financial, comm, AI safety suites + E2E golden path (\$36–37)
- ☐ 19. DevOps: Docker image, compose, IaC, blue/green, backup/restore drill (\$38–40, \$43–44)
- ☐ 20. Obs: SLOs, OTel traces, alerts with runbooks, AI/collection dashboards (\$41–42)
- ☐ 21. Monetize: 4 tiers + metered overage, usage meter, 80/100% alerts (\$45–47)
- ☐ 22. GTM: pilot referral loop, case studies, activation metric, churn watch (\$48–49, \$59)
- ☐ 23. Docs: ADR, ERD, OpenAPI, runbooks, release notes (\$53, \$57)
- ☐ 24. Gates: Definition of Done, deploy checklist, risk register review monthly (\$56, \$58, \$50)

REQUIREMENT No 'ready for pilots' claim until the money, tenant, and AI safety rows are green. Design can be rough; those three cannot.

SECTION 64

Glossary

Shared language — every term means one thing

Term	Definition
Outstanding	Total unpaid on an invoice; sum of balances per customer/org
Overdue	Outstanding where due_date < today and status not paid/cancelled
Promise to Pay	Customer's stated intent with a date/amount; recorded as PromiseToPay, not a payment
CollectionPolicy	Org-wide rules: intervals, channels, quiet hours, max attempts, escalation
CollectionCampaign	Cohort-specific steps (e.g., School Fees) built on policy defaults
CommunicationEvent	One attempted send (channel, template, status, provider_msg_id, cost)
AIConversation	Stateful AI session tied to customer+invoice, with messages/actions
AIAction	Audited tool call or decision (tool, input, output, cost, prompt_version)
Idempotency key	Client-provided unique key that makes retries safe (no double-create)
WebhookEvent	Durably stored provider callback payload before processing
Tenant isolation	Guarantee that org A never reads/writes org B's rows — enforced at DB/service/API
Grounding	Requirement that AI output amounts/dates match verified tool payloads
Quiet hours	No-contact window (e.g., 20:00–08:00 WAT) enforced before any send
Escalation	Hand-off to human via SupportTicket when AI should stop
Receipt	Immutable proof of payment (invoice, amount, ref, date, balance)
SLO	

Service-level objective (e.g., 99.5% availability) with burn-down and alert

RLS

PostgreSQL Row-Level Security — DB-enforced tenant filter

Use these terms verbatim in tickets, PRs, and docs. If a new term is introduced, add it here and link the ADR.

SECTION 65

Final Product Principles

What we will and will not do — the contract that outlives this document

CollectNaija succeeds when a business gets paid and the customer still feels respected. Every decision below is a constraint that serves that outcome.

01. Simple for the owner

One screen to know who owes what; three clicks to act. Complexity lives in policy, not in UI.

02. Respectful to the customer

Professional, calm, concise, empathetic — every message identifies the business, states the facts from verified data, offers a payment link and a human path, and respects 'stop' immediately.

03. Powerful for collection teams

Clear queues (overdue, promises due, escalations), full transcripts, and accountably assigned tickets — not an inbox of forwarded WhatsApps.

04. Reliable for financial records

Backend is the source of truth. No float math, no frontend confirmation, no duplicate payment, no reminder after verified payment — ever.

05. Safe for AI

State machine over chatbot, tools over prompts, grounding over hallucination, policy gate before every send, escalation before harm.

06. Secure by design

Tenant-isolated at DB/service/API, audited everywhere, encrypted everywhere, tested with dual-org fixtures, and pen-tested before handling real money.

07. Mobile-first, offline-aware

Designed for low-end Android and 3G; honest about offline state; queued intents with idempotency; 44px targets; no overflow.

08. AI-assisted, not AI-controlled

The owner controls policy; the customer retains preferences and protections; humans handle the hard cases. Automation compounds, but judgment stays human.

09. Scalable internationally

Currency, payment rails, comm providers, languages, and time zones are config — not hard-coded 'NGN/Africa/Lagos/English'.

010. Builder's humility

Treat pricing, channels, and vertical as hypotheses until pilots prove them. Ship weekly, listen weekly, and cut scope before cutting correctness.

Closing charge

Build the system you'd trust to message your own mother about a school fee. If it isn't respectful, verifiable, and cancellable, it isn't shippable.

What to do next (72 hours)

- Founder: approve ICP (schools vs trade wedge), sign 10 pilot targets, and lock the 4-week discovery calendar.
- Product: freeze MVP scope to sections 9 + 63 checklist, cut everything else to 'Future Idea', and draft the 8-step onboarding copy.
- Eng: scaffold Django apps + tenant middleware + invoice/payment Decimal models + webhook HMAC + gateway abstraction stub with tests.
- AI: draft the 3 prompt versions (reminder, classify, conversation) + safety blocklist + grounding validator interface, but don't wire to prod until text baseline ships.

- Docs: generate `/docs/erd.pdf` and `/docs/openapi.yaml` from code within week 2 — keep this blueprint as the source, code as the truth.

FACT

The market will teach you more in one pilot week than in a month of planning. This blueprint is a strong starting point — the pilots are the finish. Go talk to 20 businesses.